

---

# Attention Is All You Need (How to implement Transformer)

*2020. 01. 06*

*Juhyuk Lee*

*sehkmg@gmail.com*

# Table of Contents

---

- **NMT Basics and Encoder-Decoder Model**
- **Abstract View of Transformer Architecture**
- **Dive into the Transformer Architecture**

# Table of Contents

---

- **NMT Basics and Encoder-Decoder Model**
- Abstract View of Transformer Architecture
- Dive into the Transformer Architecture

# NMT Basics and Encoder-Decoder Model

---

- Neural Machine Translation
  - Dataset만 이용해서 black box model을 end-to-end로 학습시켜 번역.
  - 기본적인 Deep Learning Framework.

# NMT Basics – Dataset 준비

---

- Dataset 준비

## Source (Train)

Je suis etudiant.  
Quel mois?

## Target (Train)

I am a student.  
What month?

## Source (Valid)

Je suis medecin.  
Ce mois-ci?

## Target (Valid)

I am a doctor.  
This month?

## Source (Test)

Je suis enseignant.  
Quel jour?

## Target (Test)

I am a teacher.  
What day?

# NMT Basics – Data preprocessing

---

- Make batches - tokenize

## Source Batch (Train)

[Je, suis, etudiant, .]  
[Quel, mois, ?]

## Target Batch (Train)

[I, am, a, student, .]  
[What, month, ?]

## Source Batch (Valid)

[Je, suis, medecin, .]  
[Ce, mois, -, ci, ?]

## Target Batch (Valid)

[I, am, a, doctor, .]  
[This, month, ?]

## Source Batch (Test)

[Je, suis, enseignant, .]  
[Quel, jour, ?]

# NMT Basics – Data preprocessing

---

- Make batches – add <sos>, <eos> token

## Source Batch (Train)

[Je, suis, etudiant, .]  
[Quel, mois, ?]

## Target Batch (Train)

[<sos>, I, am, a, student, ., <eos>]  
[<sos>, What, month, ?, <eos>]

## Source Batch (Valid)

[Je, suis, medecin, .]  
[Ce, mois, -, ci, ?]

## Target Batch (Valid)

[<sos>, I, am, a, doctor, ., <eos>]  
[<sos>, This, month, ?, <eos>]

## Source Batch (Test)

[Je, suis, enseignant, .]  
[Quel, jour, ?]

# NMT Basics – Data preprocessing

---

- Make batches – padding

## Source Batch (Train)

[Je, suis, etudiant, .]  
[Quel, mois, ?, <pad>]

## Target Batch (Train)

[<sos>, I, am, a, student, ., <eos>]  
[<sos>, What, month, ?, <eos>, <pad>, <pad>]

## Source Batch (Valid)

[Je, suis, medecin, ., <pad>]  
[Ce, mois, -, ci, ?]

## Target Batch (Valid)

[<sos>, I, am, a, doctor, ., <eos>]  
[<sos>, This, month, ?, <eos>, <pad>, <pad>]

## Source Batch (Test)

[Je, suis, enseignant, .]  
[Quel, jour, ?, <pad>]



# NMT Basics – Data preprocessing

---

- Numericalize – make vocabulary from train dataset

## Source (Train)

Je suis etudiant.  
Quel mois?

## Target (Train)

I am a student.  
What month?

## Source Vocab

0: <sos>  
1: <eos>  
2: <pad>  
3: <unk>  
4: Je  
5: suis  
6: etudiant  
7: .  
8: Quel  
9: mois  
10: ?  
11: .....

## Target Vocab

0: <sos>  
1: <eos>  
2: <pad>  
3: <unk>  
4: I  
5: am  
6: a  
7: student  
8: .  
9: What  
10: month  
11: ?  
12: .....

# NMT Basics – Data preprocessing

---

- Numericalize – handle out-of-vocabulary words

## Source Batch (Train)

[Je, suis, etudiant, .]  
[Quel, mois, ?, <pad>]

## Target Batch (Train)

[<sos>, I, am, a, student, ., <eos>]  
[<sos>, What, month, ?, <eos>, <pad>, <pad>]

## Source Batch (Valid)

[Je, suis, <unk>, ., <pad>]  
[<unk>, mois, <unk>, <unk>, ?]

## Target Batch (Valid)

[<sos>, I, am, a, <unk>, ., <eos>]  
[<sos>, <unk>, month, ?, <eos>, <pad>, <pad>]

## Source Batch (Test)

[Je, suis, <unk>, .]  
[Quel, <unk>, ?, <pad>]

# NMT Basics – Data preprocessing

---

- Numericalize – numericalize

## Source Batch (Train)

[4, 5, 6, 7]  
[8, 9, 10, 2]

## Target Batch (Train)

[0, 4, 5, 6, 7, 8, 1]  
[0, 9, 10, 11, 1, 2, 2]

## Source Batch (Valid)

[4, 5, 3, 7, 2]  
[3, 9, 3, 3, 10]

## Target Batch (Valid)

[0, 4, 5, 6, 3, 8, 1]  
[0, 3, 10, 11, 1, 2, 2]

## Source Batch (Test)

[4, 5, 3, 7]  
[8, 3, 10, 2]


# Encoder-Decoder Model - Training

Source Batch (Train)

4	5	6	7
8	9	10	2

Target Batch (Train)

0	4	5	6	7	8	1
0	9	10	11	1	2	2



Target Batch (Input)

0	4	5	6	7	8	1
0	9	10	11	1	2	2

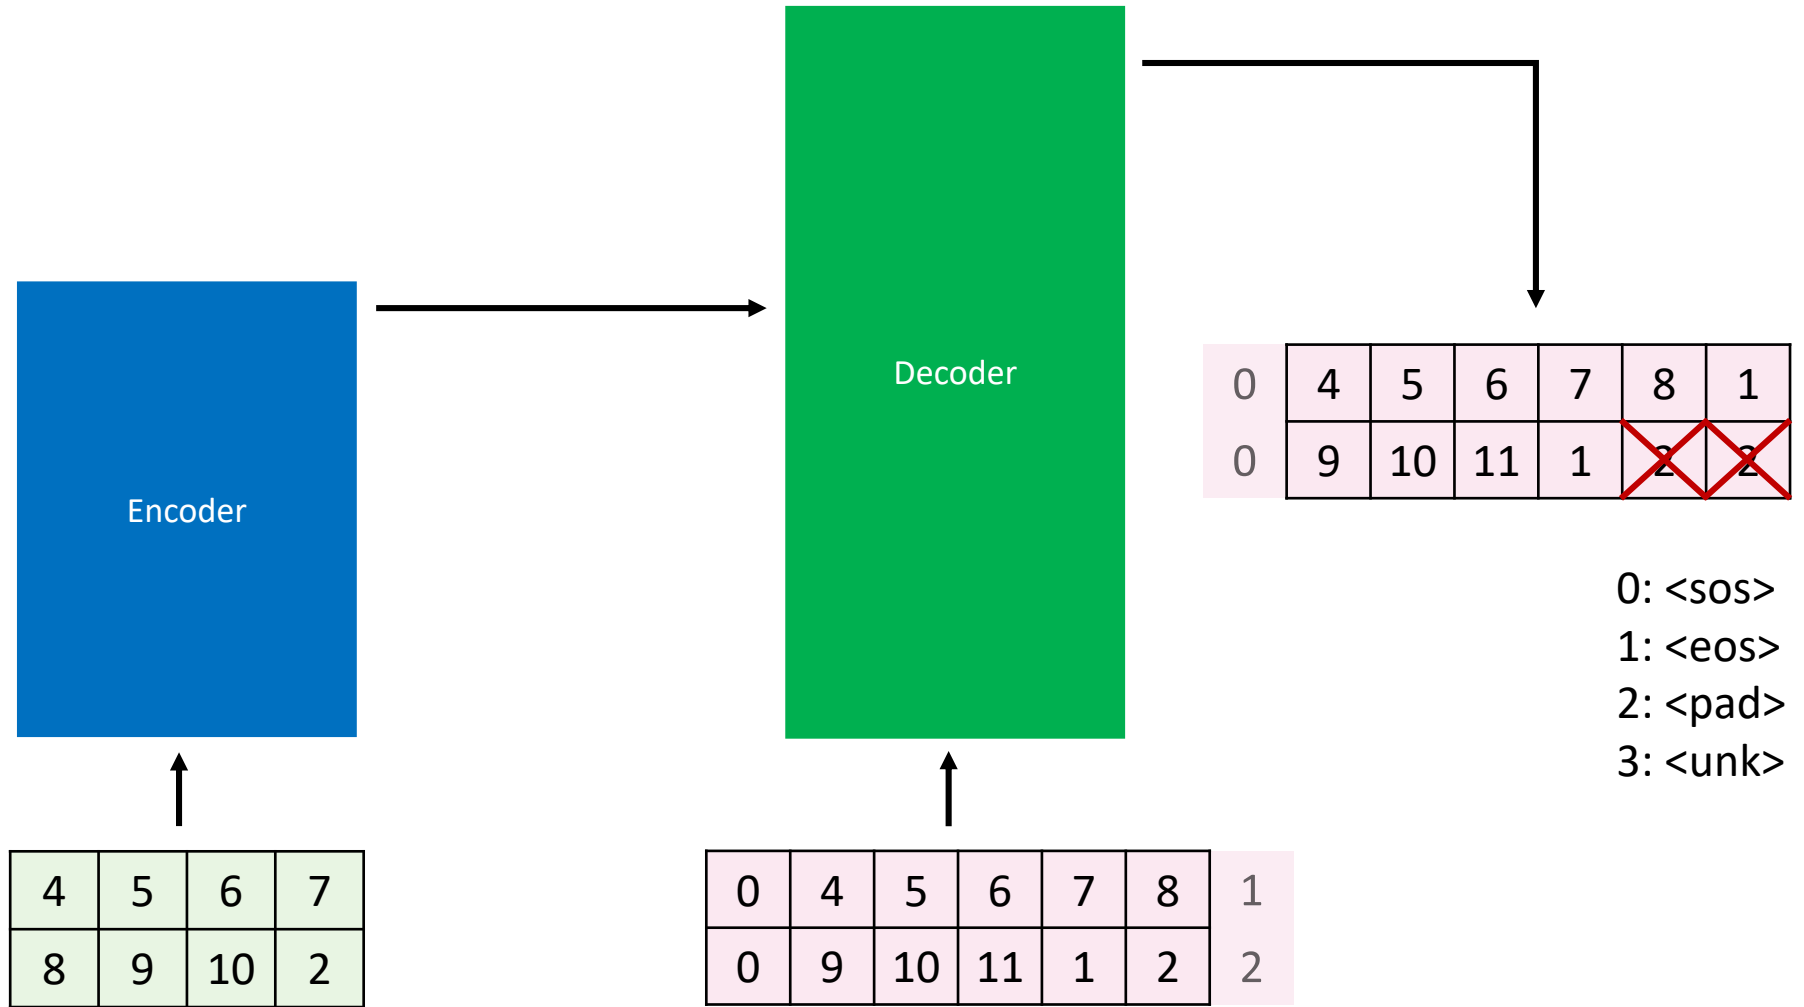
Target Batch (True Output)

0	4	5	6	7	8	1
0	9	10	11	1	2	2

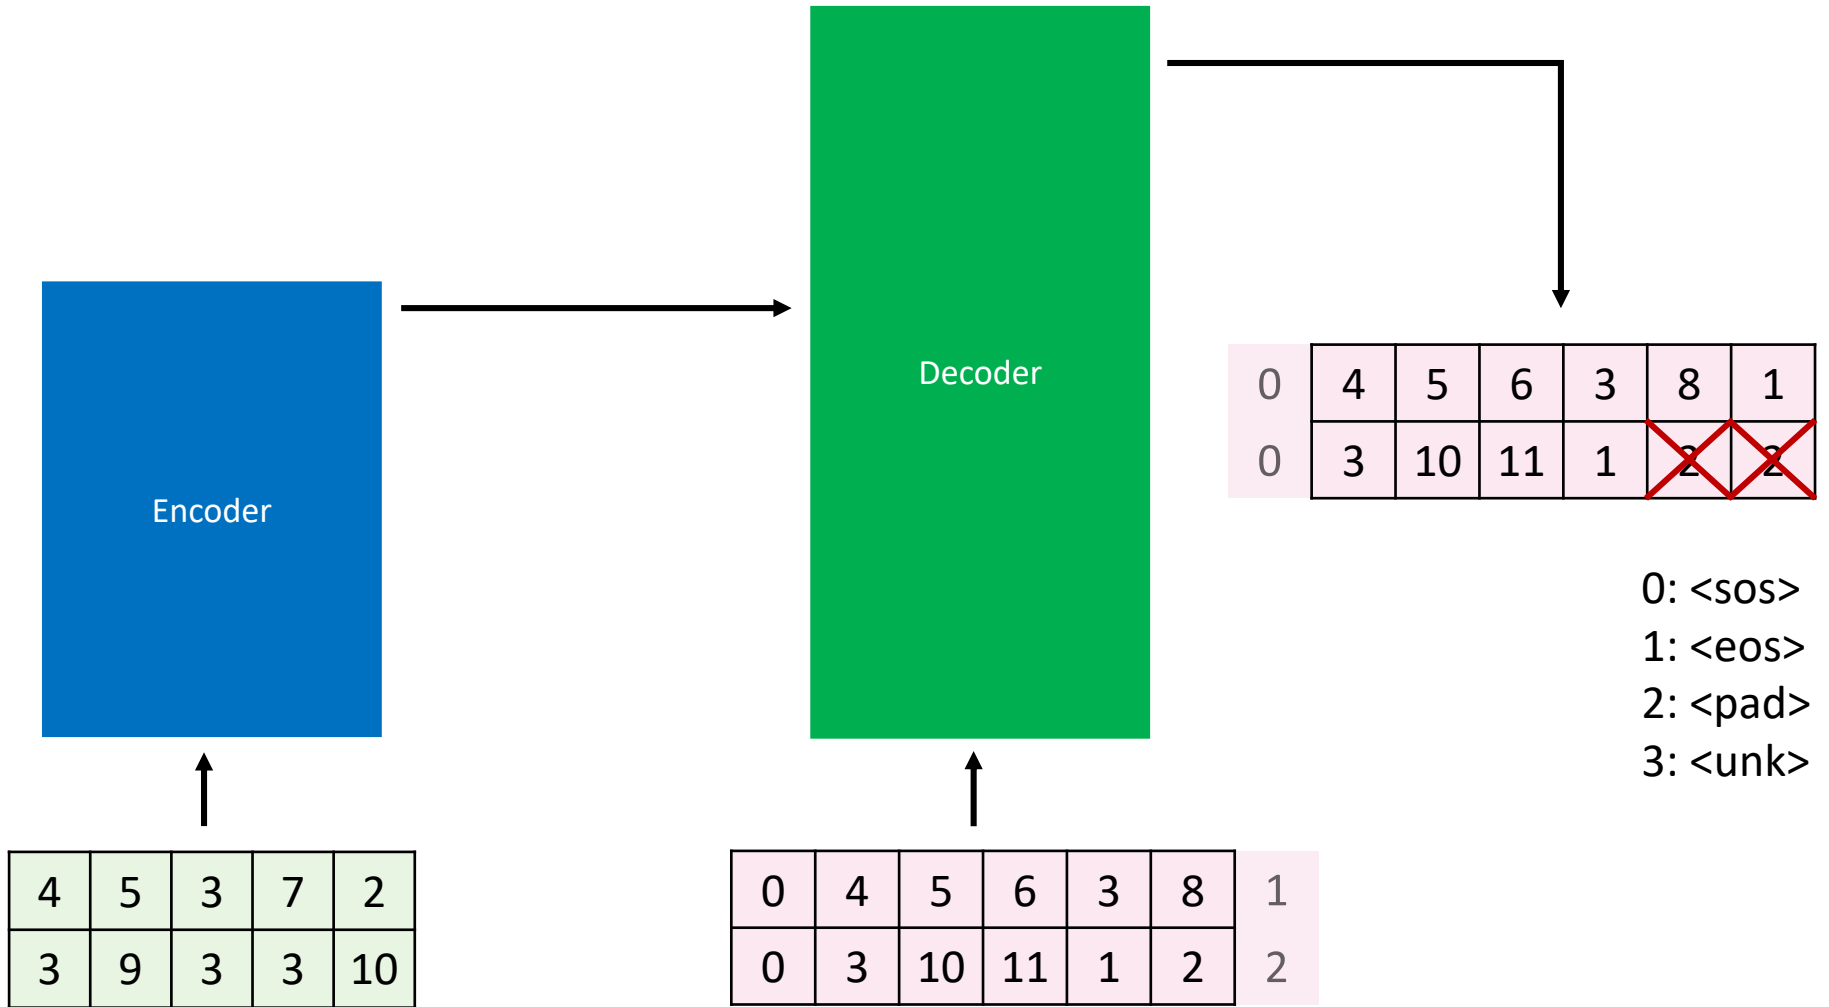
Encoder

Decoder


# Encoder-Decoder Model - Training



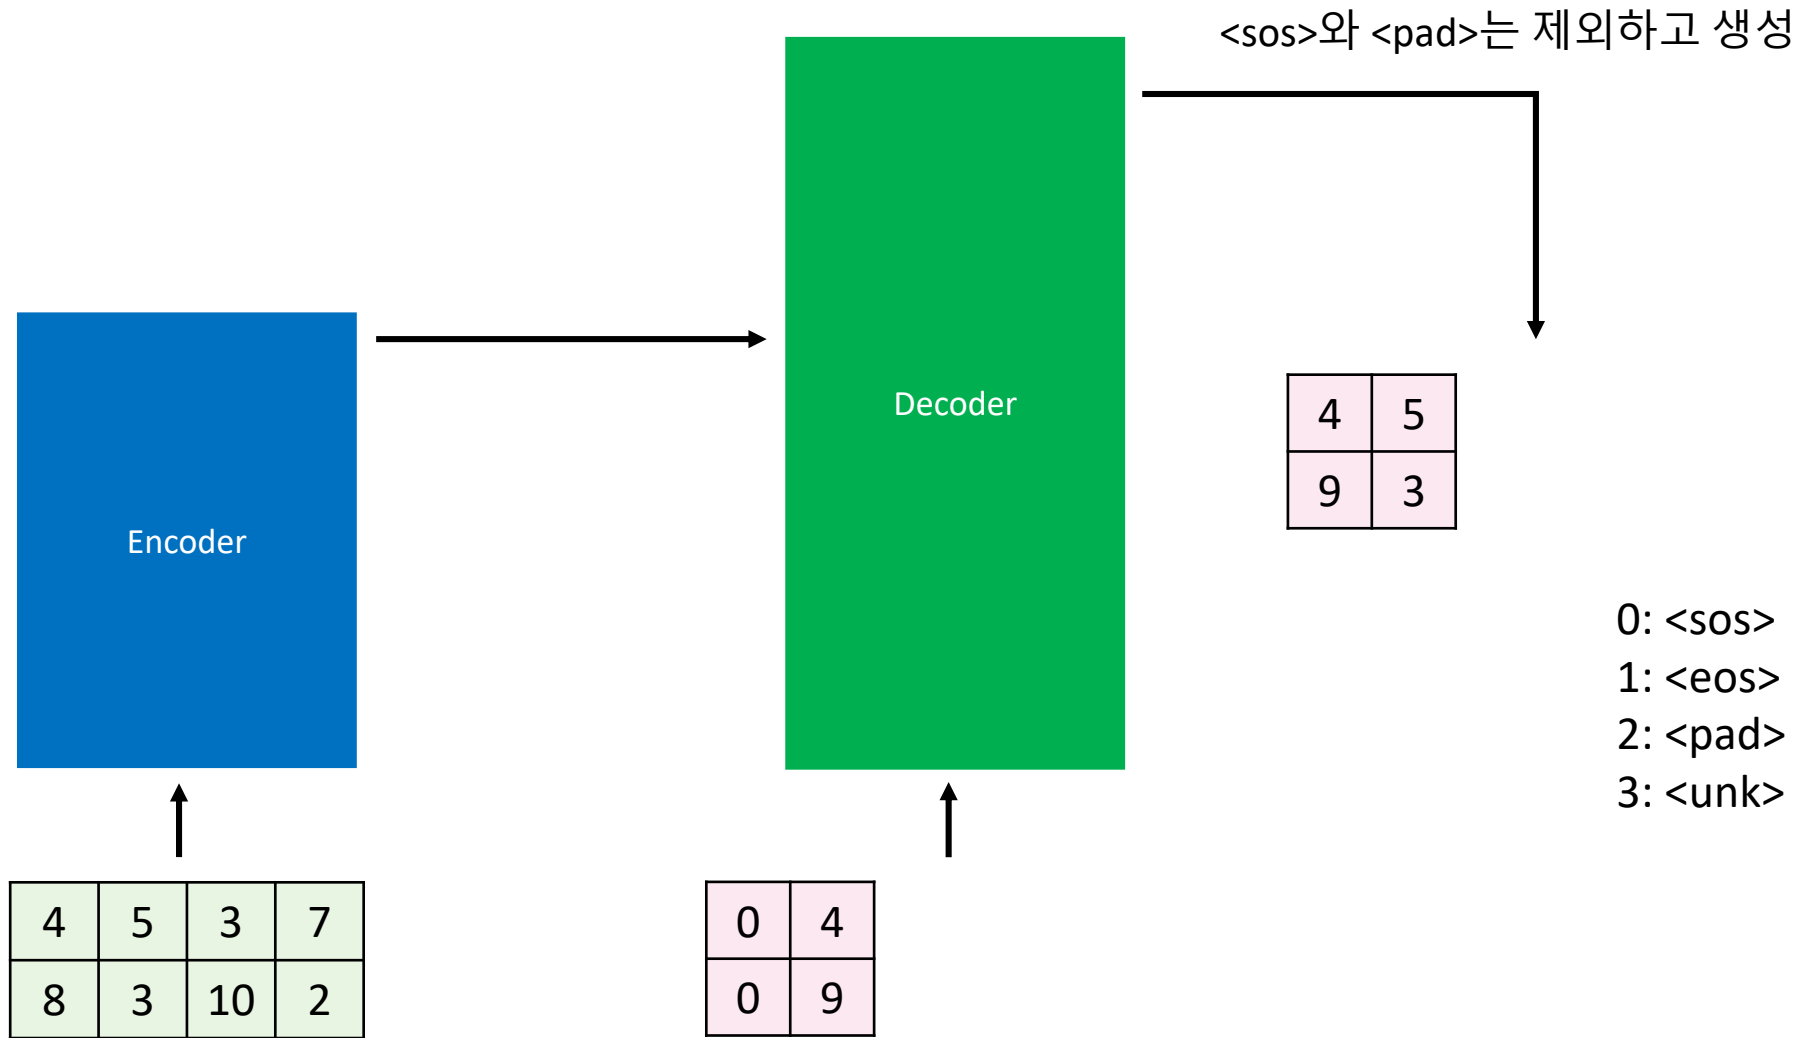
# Encoder-Decoder Model - Validation



# Encoder-Decoder Model – Test (Inference)

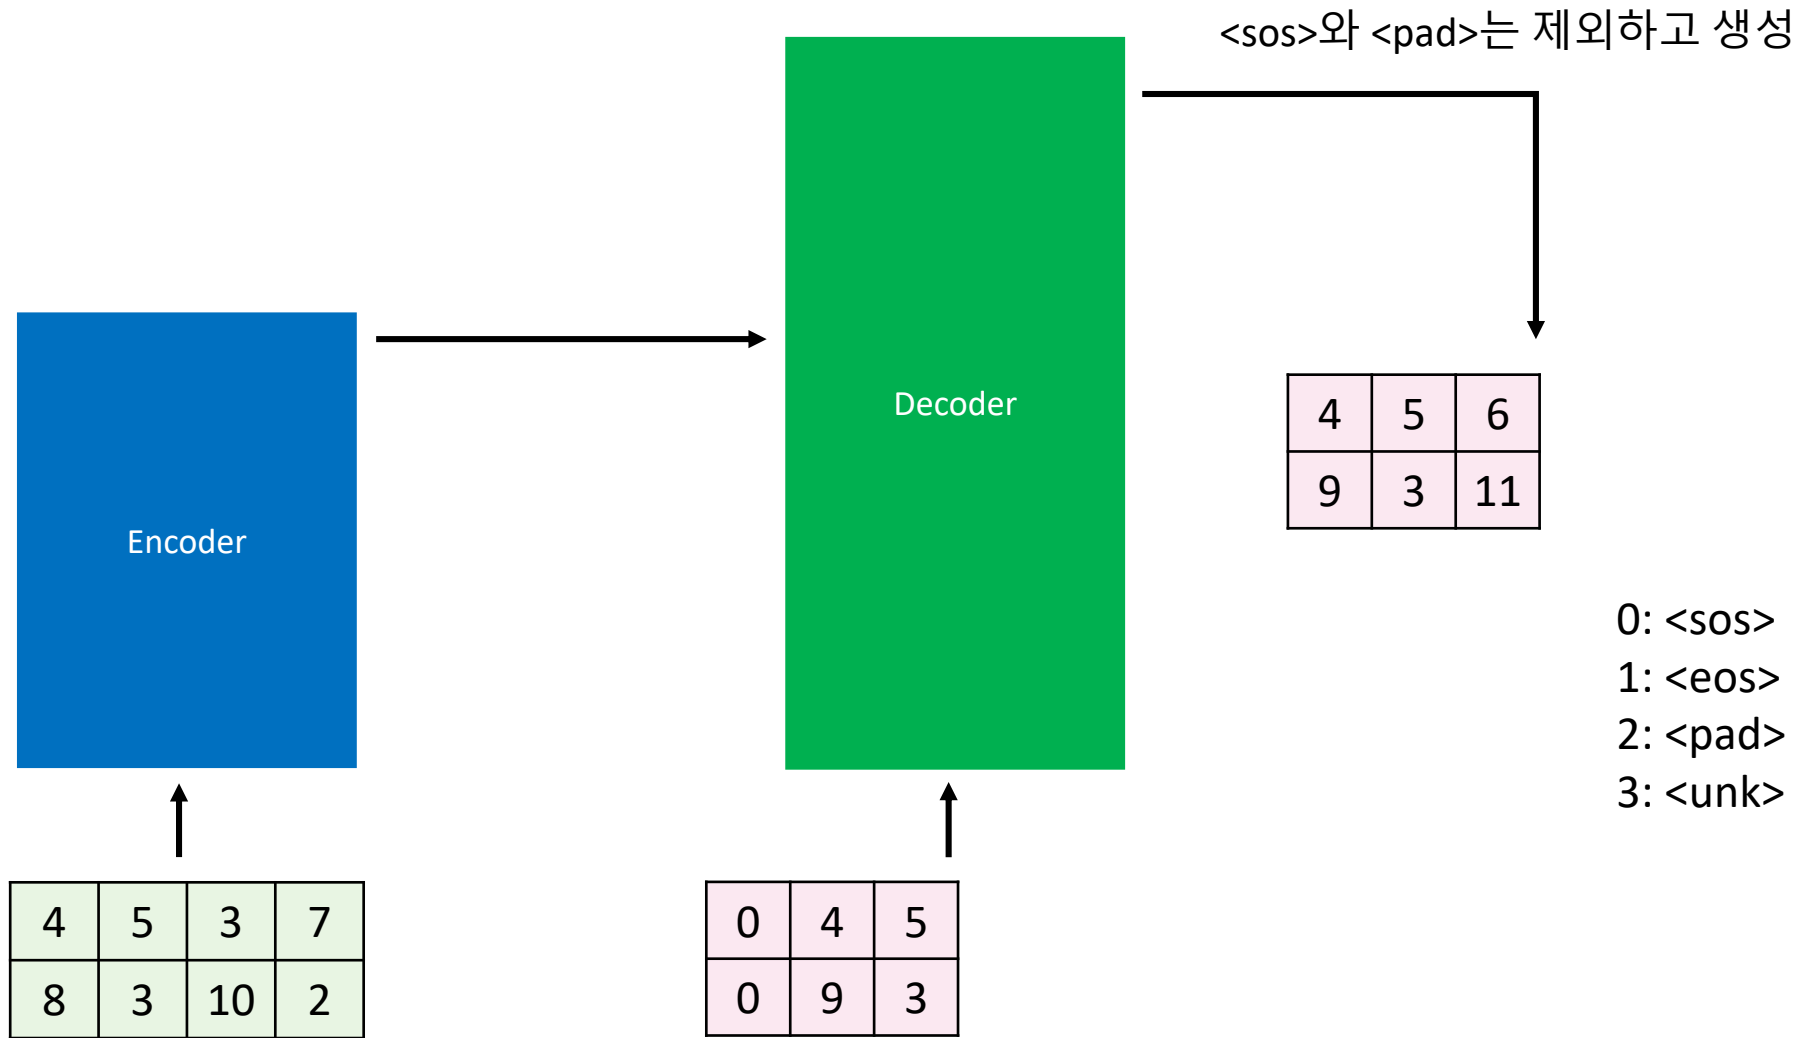


# Encoder-Decoder Model – Test (Inference)

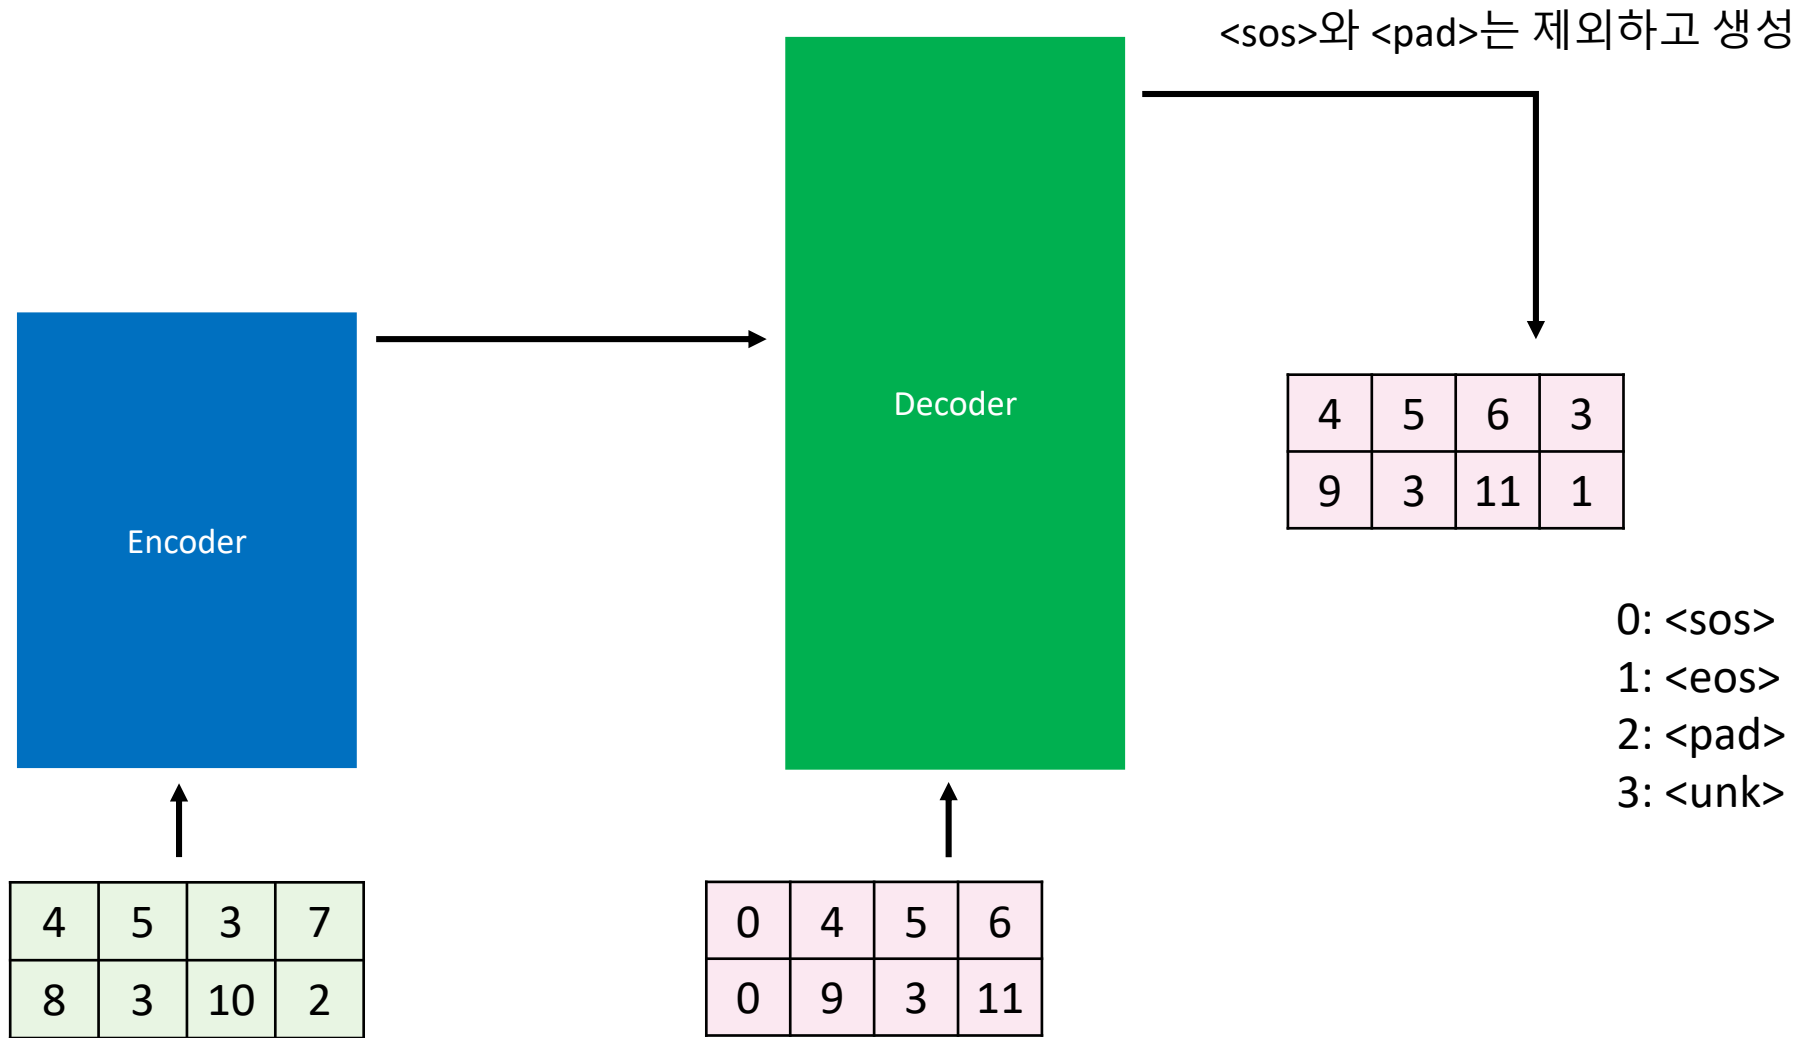




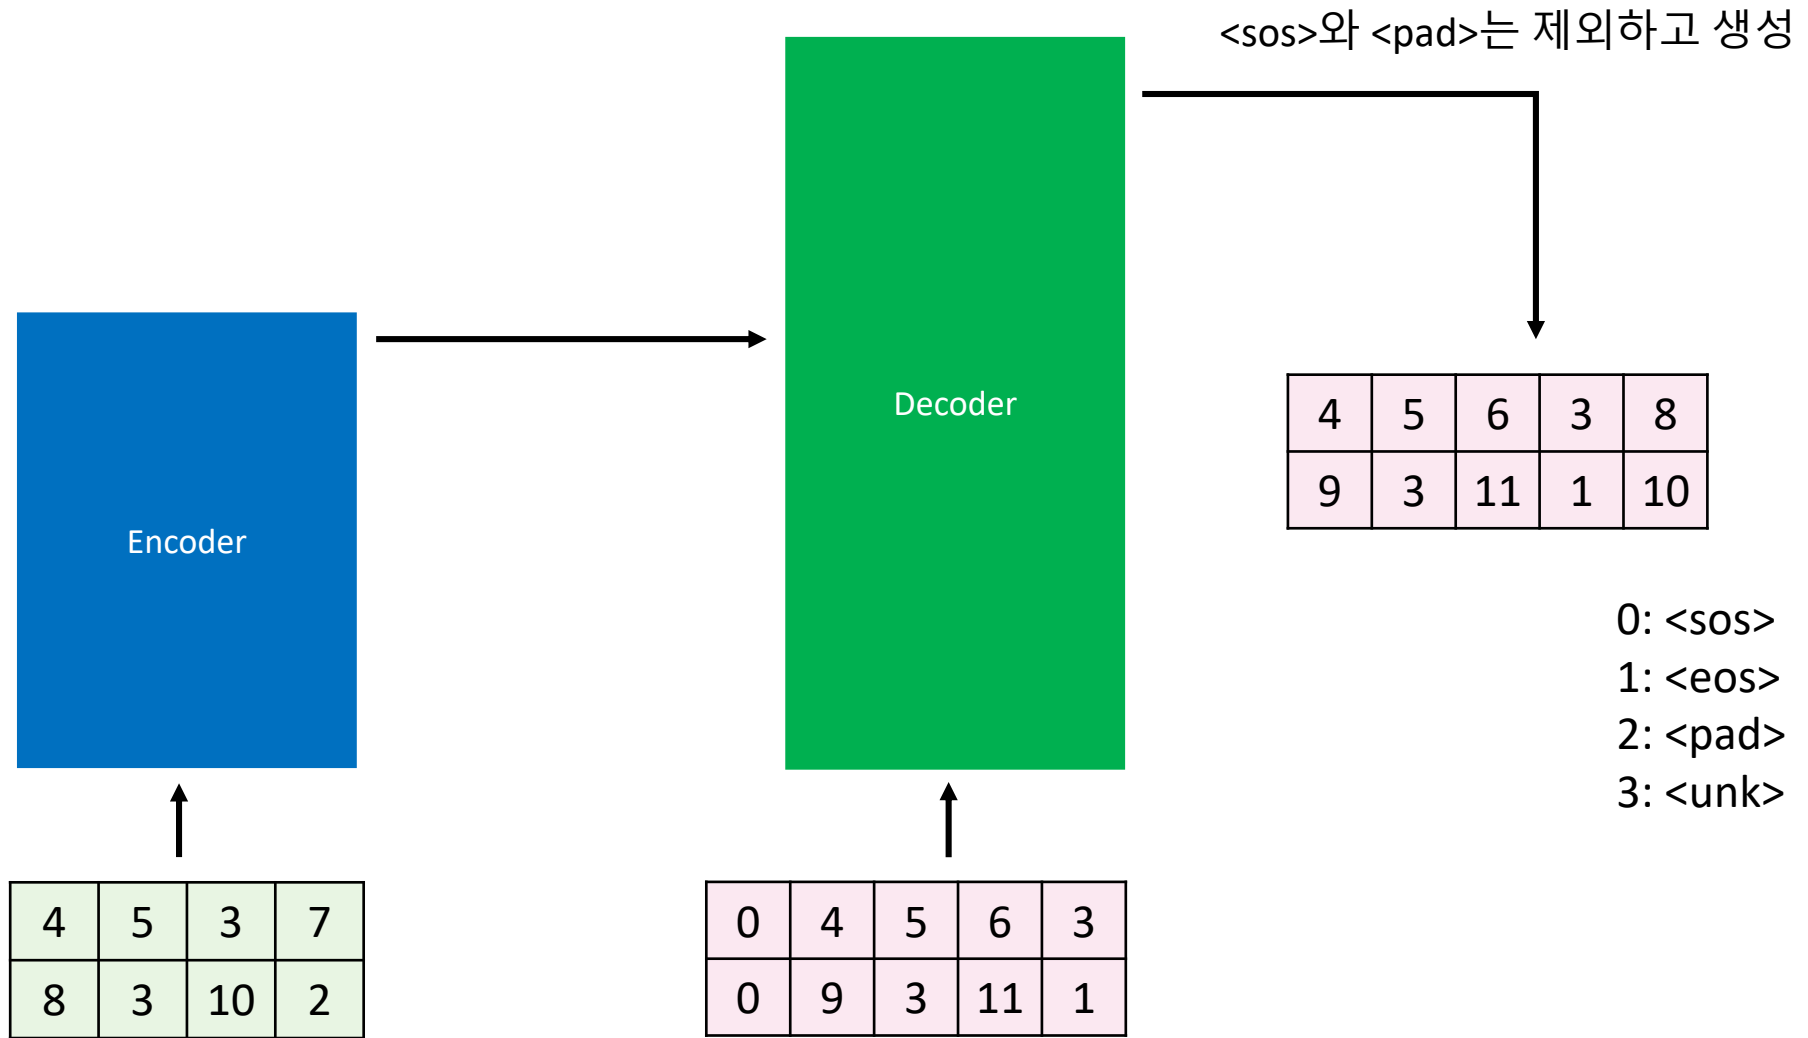
# Encoder-Decoder Model – Test (Inference)



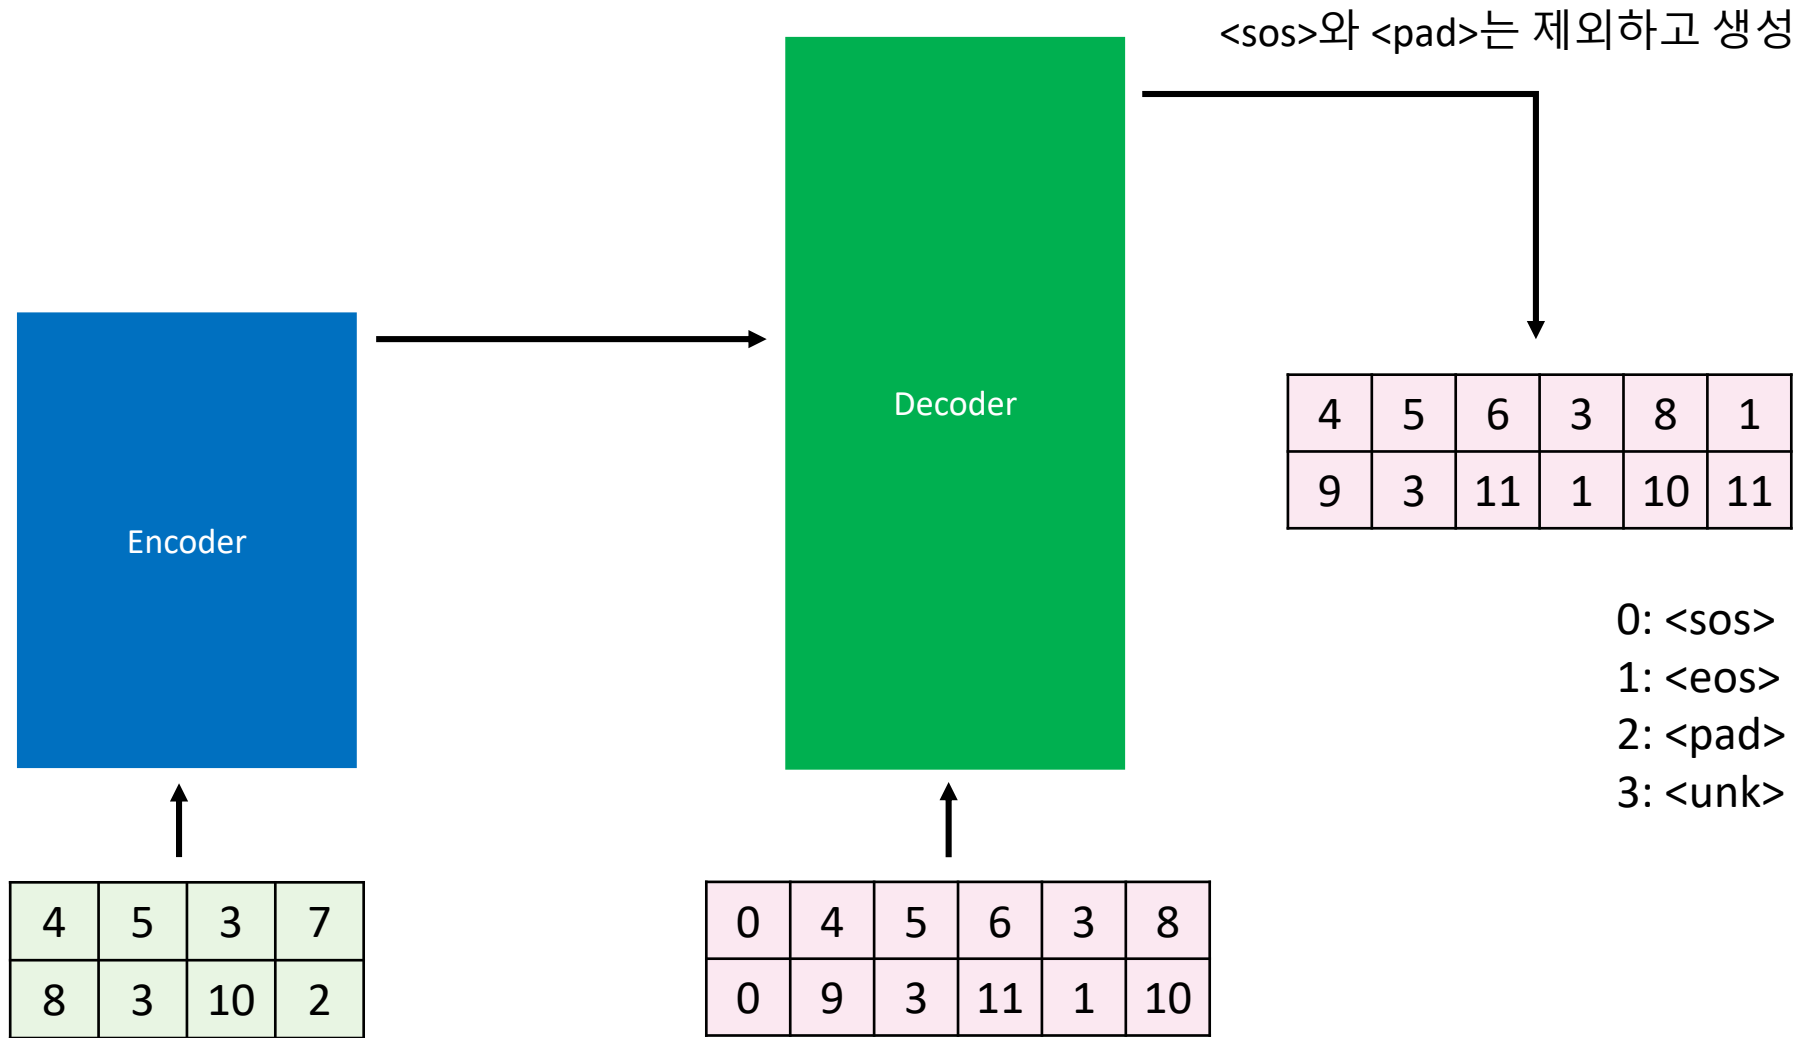
# Encoder-Decoder Model – Test (Inference)



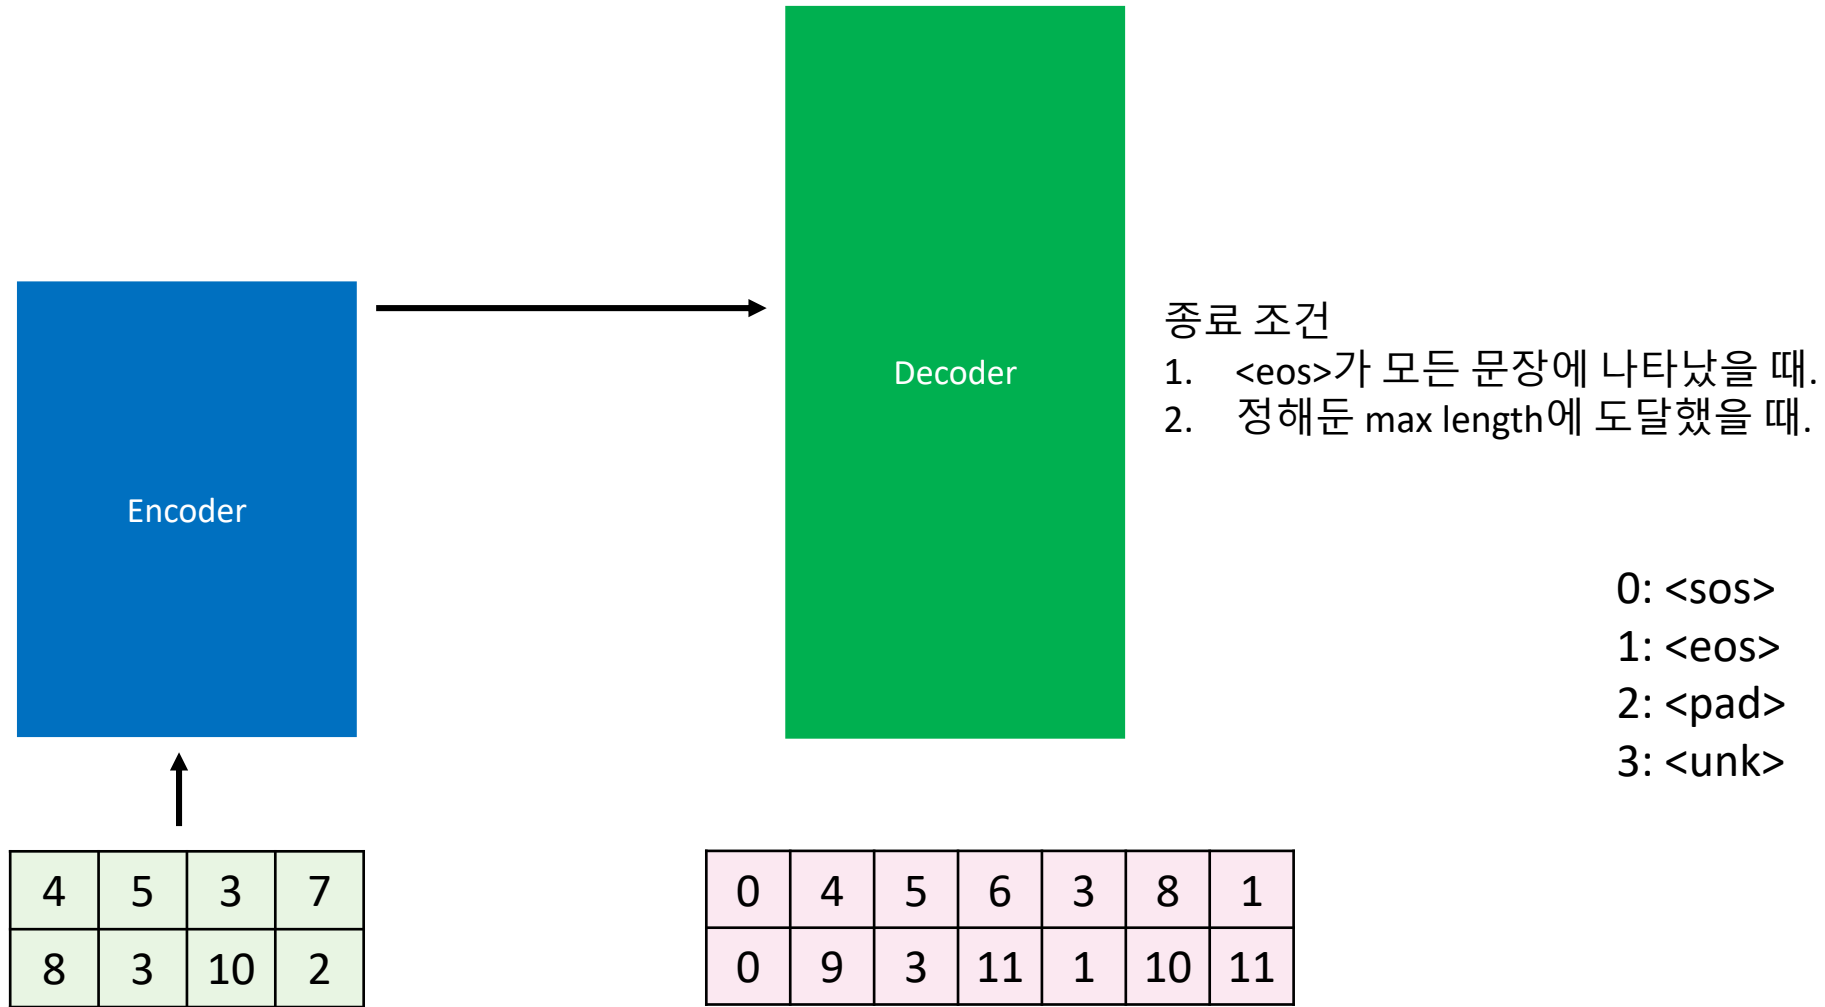
# Encoder-Decoder Model – Test (Inference)



# Encoder-Decoder Model – Test (Inference)



# Encoder-Decoder Model – Test (Inference)



# NMT Basics – Compute BLEU Score

## Target Vocab

0: <sos>  
1: <eos>  
2: <pad>  
3: <unk>  
4: I  
5: am  
6: a  
7: student  
8: .  
9: What  
10: month  
11: ?  
12: .....

0	4	5	6	3	8	1
0	9	3	11	1	10	11



0	4	5	6	3	8	1
0	9	3	11	1	10	11



[I, am, a, <unk>, .]  
[What, <unk>, ?]



I am a <unk>.  
What <unk>?

BLEU score



## Target (Test)

I am a teacher.  
What day?

# NMT Basics and Encoder-Decoder Model

---

- **NMT Basics**

- Dataset 준비
- Data preprocessing
  - Make batches
    - tokenize
    - add <sos>, <eos> token
    - padding
  - Numericalize
    - make vocabulary from train dataset
    - handle out-of-vocabulary words
    - numericalize
- Compute BLEU Score

- **Encoder-Decoder Model**

- Training
- Validation
- Test (Inference)

# NMT Basics and Encoder-Decoder Model

- **NMT Basics (Done!)**

- Dataset 준비: Multi30k English to German Translation Dataset
- Data preprocessing
  - Make batches
    - tokenize
    - add <sos>, <eos> token
    - padding
  - Numericalize
    - make vocabulary from train dataset
    - handle out-of-vocabulary words
    - numericalize
- Compute BLEU Score

- **Encoder-Decoder Model**

- Training
- Validation
- Test (Inference)

[https://github.com/sehkmng/NMT\\_practice](https://github.com/sehkmng/NMT_practice)

```
# TODO: train
for epoch in range(args.epochs):
    for src_batch, tgt_batch in train_loader:
        pass

    # TODO: validation
    for src_batch, tgt_batch in valid_loader:
        pass
```

```
for src_batch, tgt_batch in test_loader:
    # TODO: predict pred_batch from src_b
    pred_batch = tgt_batch

    # every sentences in pred_batch should
    # every <pad> token (index: 2) should
    # example of pred_batch:
    # [[0, 5, 6, 7, 1],
    #  [0, 4, 9, 1, 2],
    #  [0, 6, 1, 2, 2]]
```

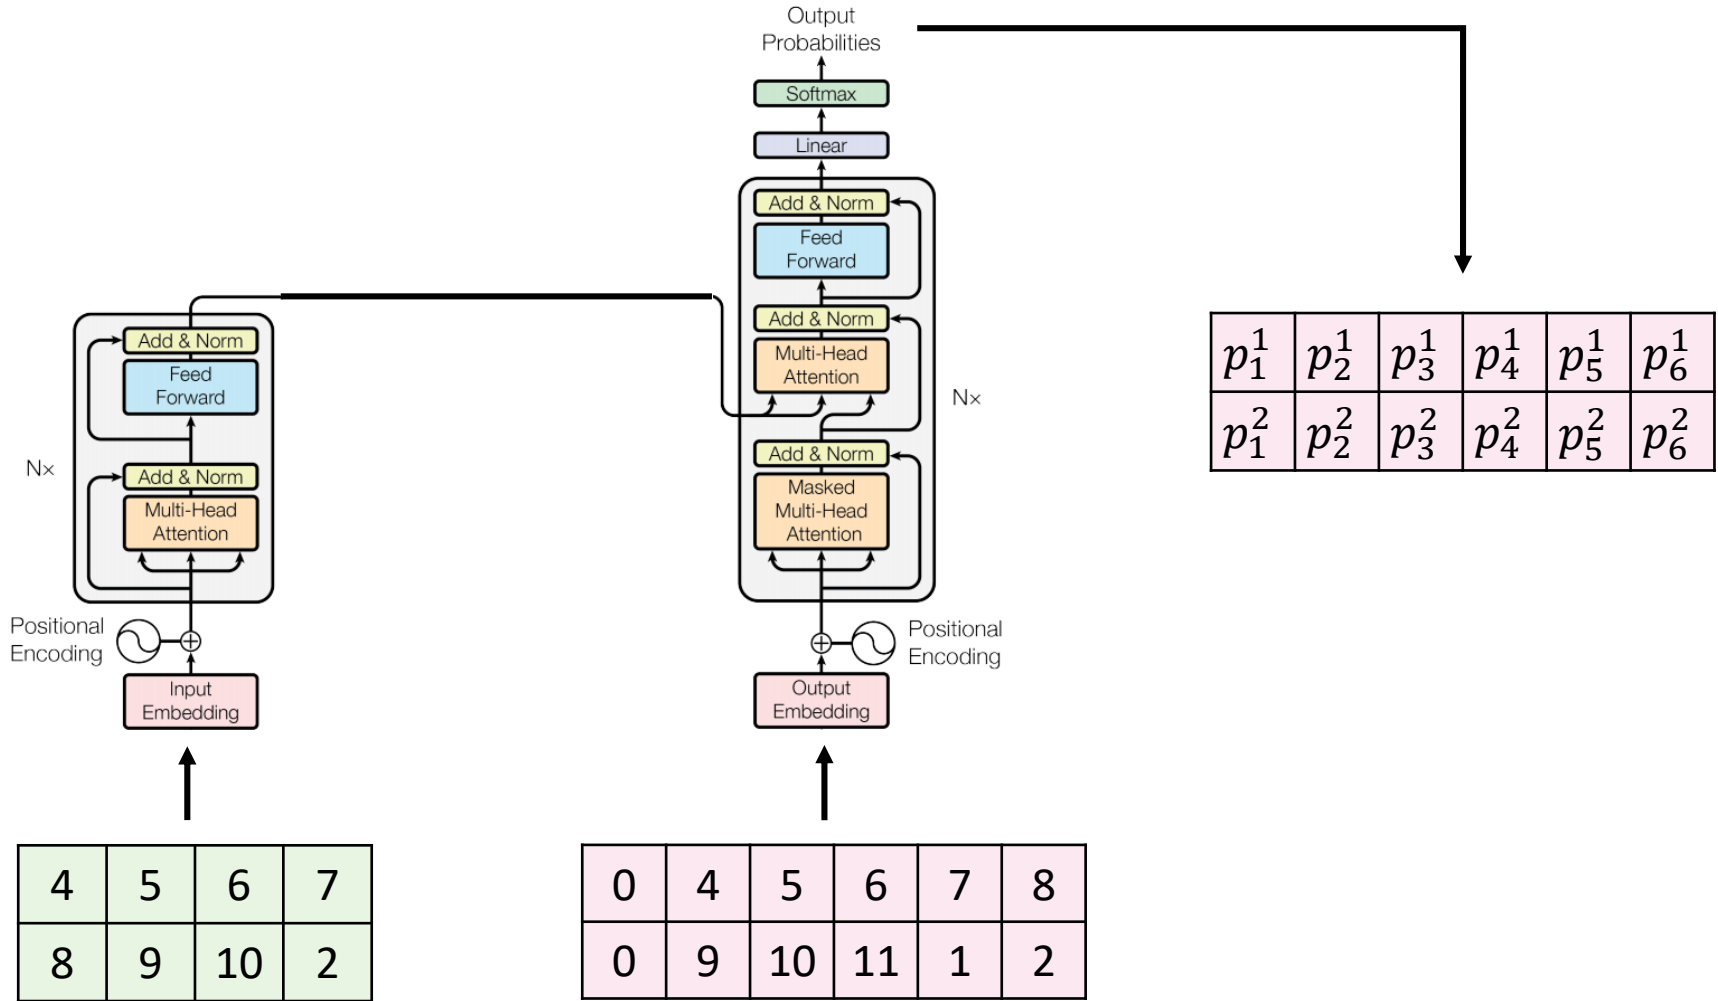


# Table of Contents

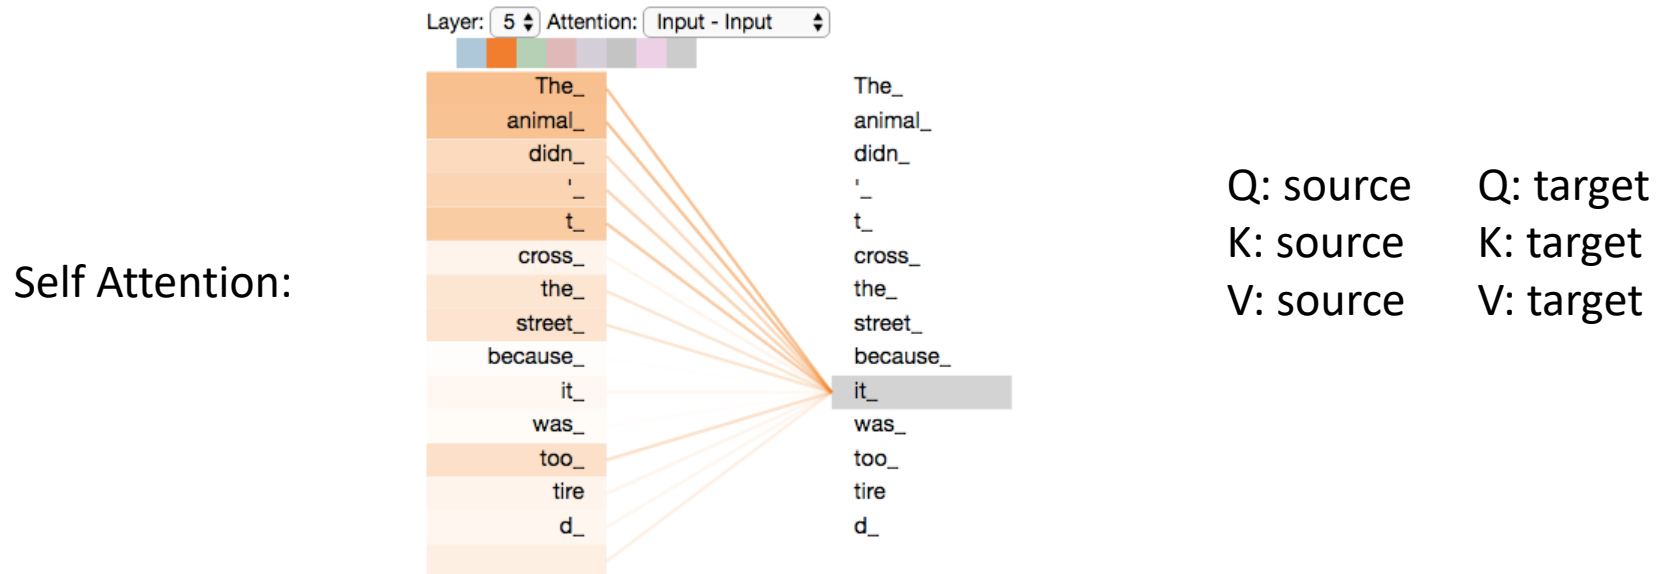

---

- NMT Basics and Encoder-Decoder Model
- **Abstract View of Transformer Architecture**
- Dive into the Transformer Architecture

# Abstract View of Transformer Architecture




# Self Attention



# Multi-Head Attention


Original Attention: Je suis étudiant . <sos> I am a student .




Attention for student:  $\alpha_1 \times \text{Je} + \alpha_2 \times \text{suis} + \alpha_3 \times \text{étudiant}$

Q: target  
K: source  
V: source


Scaled Dot-Product Attention



Multi-Head Attention



# Abstract View of Transformer Architecture



# Table of Contents

---


- NMT Basics and Encoder-Decoder Model
- Abstract View of Transformer Architecture
- **Dive into the Transformer Architecture**

# Attention is a Weighted Sum


- Query: 주인공 문장.



- Key, Value: 주인공 문장이 보는 문장.



- Query는 Key와 연산하여 weight를 구한다.



# Attention is a Weighted Sum

- $d_k$ 로 나눠주고 원치 않는 정보를 masking을 통해 지운 후 Softmax를 취한다.

$$\text{Softmax}\left(\frac{w_{11}}{d_k}, \frac{\cancel{w_{12}}}{\cancel{d_k}}, \frac{w_{13}}{d_k}\right) = (\alpha_{11}, 0, \alpha_{13})$$

$$\text{Softmax}\left(\frac{w_{21}}{d_k}, \frac{w_{22}}{d_k}, \frac{\cancel{w_{23}}}{\cancel{d_k}}\right) = (\alpha_{21}, \alpha_{22}, 0)$$

- Value를 대상으로 Weighted Sum을 한다.

$$\alpha_{11} \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array} + 0 \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array} + \alpha_{13} \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array}$$

$$\alpha_{21} \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array} + \alpha_{22} \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array} + 0 \cdot \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \end{array}$$

- 최종 output.






# Attention in Matrix Form

- Query: 주인공 문장.


- Key, Value: 주인공 문장이 보는 문장.


- Query는 Key와 연산하여 weight를 구한다.



# Attention in Matrix Form

- $d_k$ 로 나눠주고 원치 않는 정보를 masking을 통해 지운 후 Softmax를 취한다.

$$\text{Softmax}\left(\frac{\begin{array}{|c|c|c|} \hline w_{11} & \cancel{w_{12}} & w_{13} \\ \hline w_{21} & w_{22} & \cancel{w_{23}} \\ \hline \end{array}}{d_k}\right) \longrightarrow \begin{array}{|c|c|c|} \hline \alpha_{11} & 0 & \alpha_{13} \\ \hline \alpha_{21} & \alpha_{22} & 0 \\ \hline \end{array}$$

- Value를 대상으로 Weighted Sum을 한다.


$$\begin{array}{|c|c|c|} \hline \alpha_{11} & 0 & \alpha_{13} \\ \hline \alpha_{21} & \alpha_{22} & 0 \\ \hline \end{array} \times \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \text{medium green} & \text{medium green} & \text{medium green} & \text{medium green} \\ \hline \text{dark green} & \text{dark green} & \text{dark green} & \text{dark green} \\ \hline \end{array}$$

- 최종 output.



light blue	light blue	light blue	light blue
dark blue	dark blue	dark blue	dark blue

# Attention in Implementation

- Query: 주인공 문장.




- Key, Value: 주인공 문장이 보는 문장.





# Multi-Head Attention

- Query




- Key, Value




# Multi-Head Attention

- Concatenate Outputs


Out1 Out2 Out3



- Match the dimension



×




=



# Multi-Head Attention

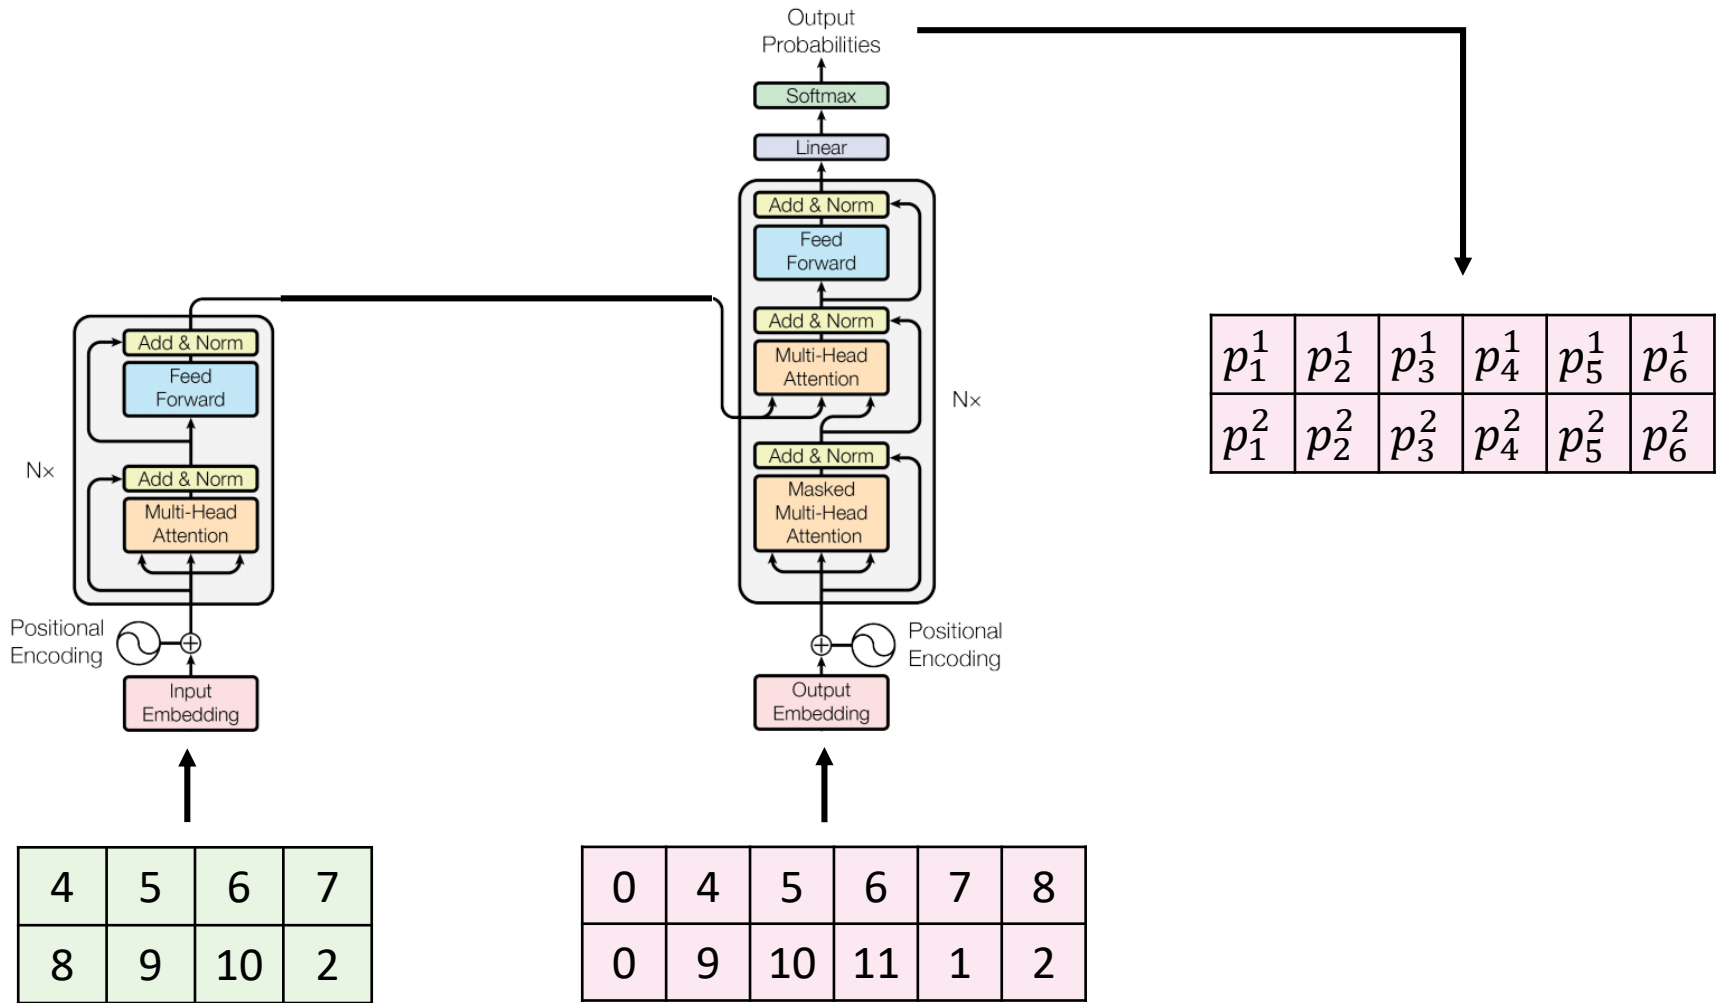
- Multi\_Head\_Attention(Query, Key, Value, Mask)

Multi\_Head\_Attention (

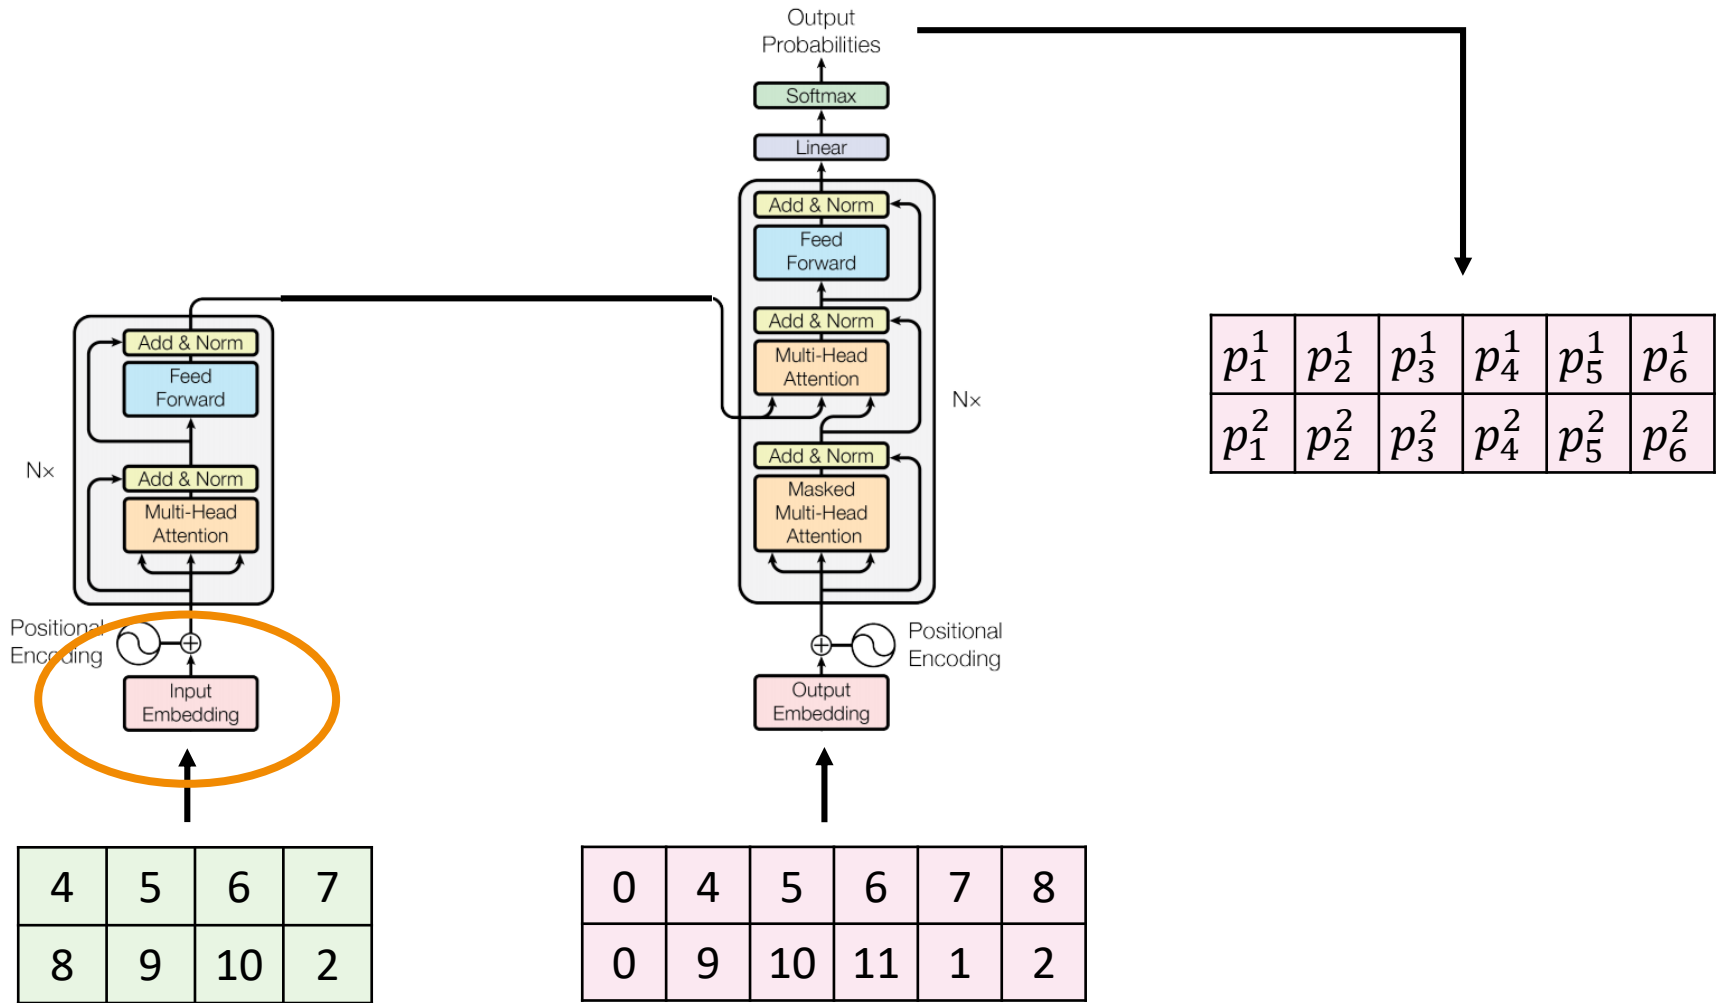


- 위 example에서는 단어 벡터의 차원이 4, Head가 3개, 각 Head의 차원이 3.
- 보통은 단어 벡터의 차원이 512, Head가 8개, 각 Head의 차원이 64.

# Dive into the Transformer Architecture



# Input Embedding





# Input Embedding

4	5	6	7
8	9	10	2

one-hot vector  
→

				1							
					1						
						1					
							1				

								1			
									1		
										1	
		1									

# Input Embedding

				1							
					1						
						1					
							1				

								1			
									1		
										1	
		1									

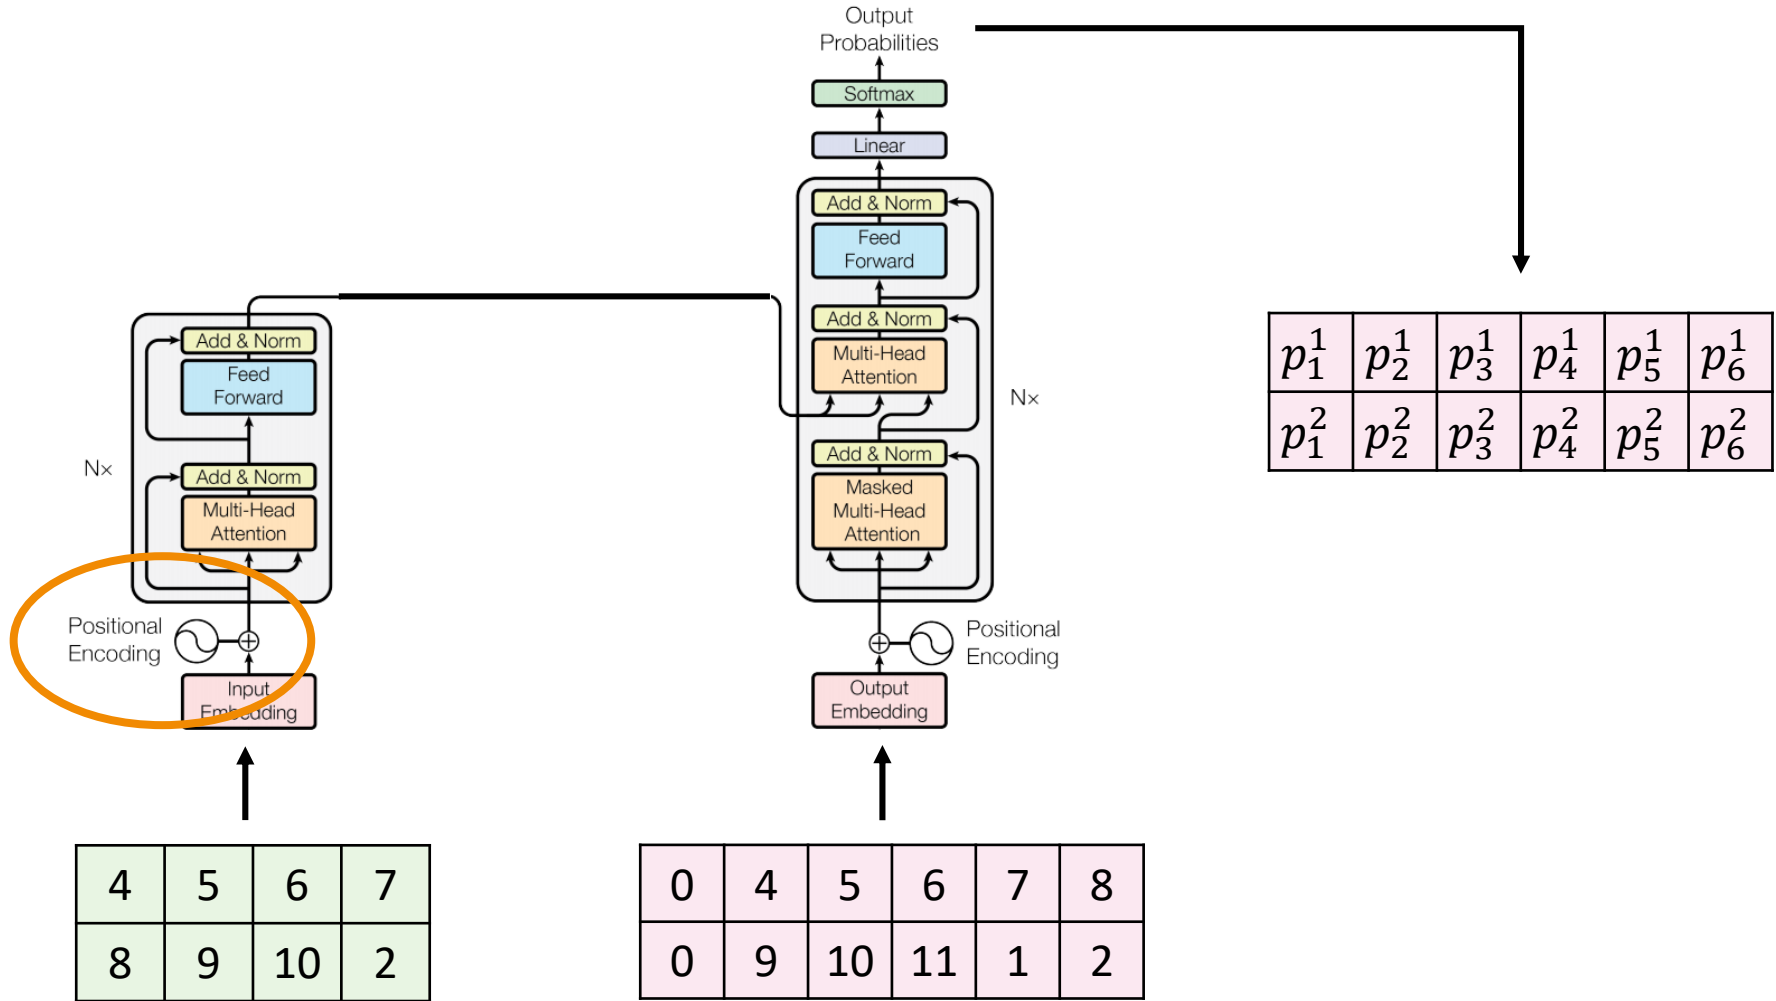
Linear map






dim=512





# Positional Encoding




# Positional Encoding

word	embedding	Positional Encoding
Je		$+$ $\cos(1) \sin\left(\frac{1}{10}\right) \cos\left(\frac{1}{10}\right) \sin\left(\frac{1}{10^2}\right) \cos\left(\frac{1}{10^2}\right) \sin\left(\frac{1}{10^3}\right) \cos\left(\frac{1}{10^3}\right) \sin\left(\frac{1}{10^4}\right)$
suis		$+$ $\cos(2) \sin\left(\frac{2}{10}\right) \cos\left(\frac{2}{10}\right) \sin\left(\frac{2}{10^2}\right) \cos\left(\frac{2}{10^2}\right) \sin\left(\frac{2}{10^3}\right) \cos\left(\frac{2}{10^3}\right) \sin\left(\frac{2}{10^4}\right)$
étudiant		$+$ $\cos(3) \sin\left(\frac{3}{10}\right) \cos\left(\frac{3}{10}\right) \sin\left(\frac{3}{10^2}\right) \cos\left(\frac{3}{10^2}\right) \sin\left(\frac{3}{10^3}\right) \cos\left(\frac{3}{10^3}\right) \sin\left(\frac{3}{10^4}\right)$

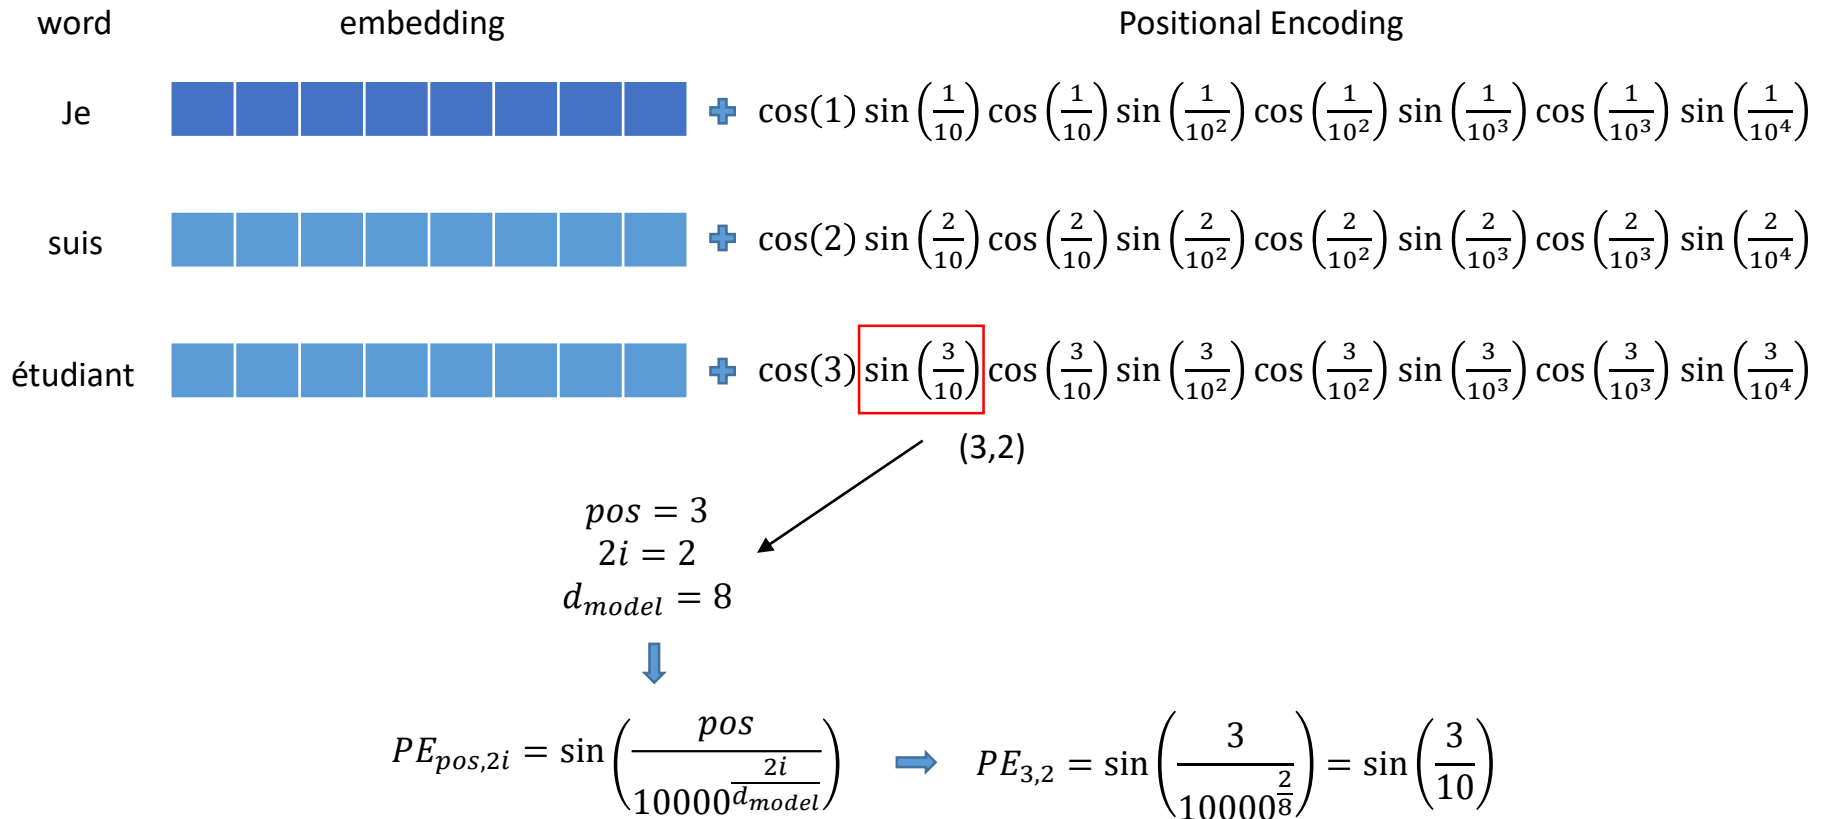
# Positional Encoding

word	embedding	Positional Encoding
Je		$\oplus \cos(1) \sin\left(\frac{1}{10}\right) \cos\left(\frac{1}{10}\right) \sin\left(\frac{1}{10^2}\right) \cos\left(\frac{1}{10^2}\right) \sin\left(\frac{1}{10^3}\right) \cos\left(\frac{1}{10^3}\right) \sin\left(\frac{1}{10^4}\right)$
suis		$\oplus \cos(2) \sin\left(\frac{2}{10}\right) \cos\left(\frac{2}{10}\right) \sin\left(\frac{2}{10^2}\right) \cos\left(\frac{2}{10^2}\right) \sin\left(\frac{2}{10^3}\right) \cos\left(\frac{2}{10^3}\right) \sin\left(\frac{2}{10^4}\right)$
étudiant		$\oplus \cos(3) \sin\left(\frac{3}{10}\right) \cos\left(\frac{3}{10}\right) \sin\left(\frac{3}{10^2}\right) \cos\left(\frac{3}{10^2}\right) \sin\left(\frac{3}{10^3}\right) \cos\left(\frac{3}{10^3}\right) \sin\left(\frac{3}{10^4}\right)$

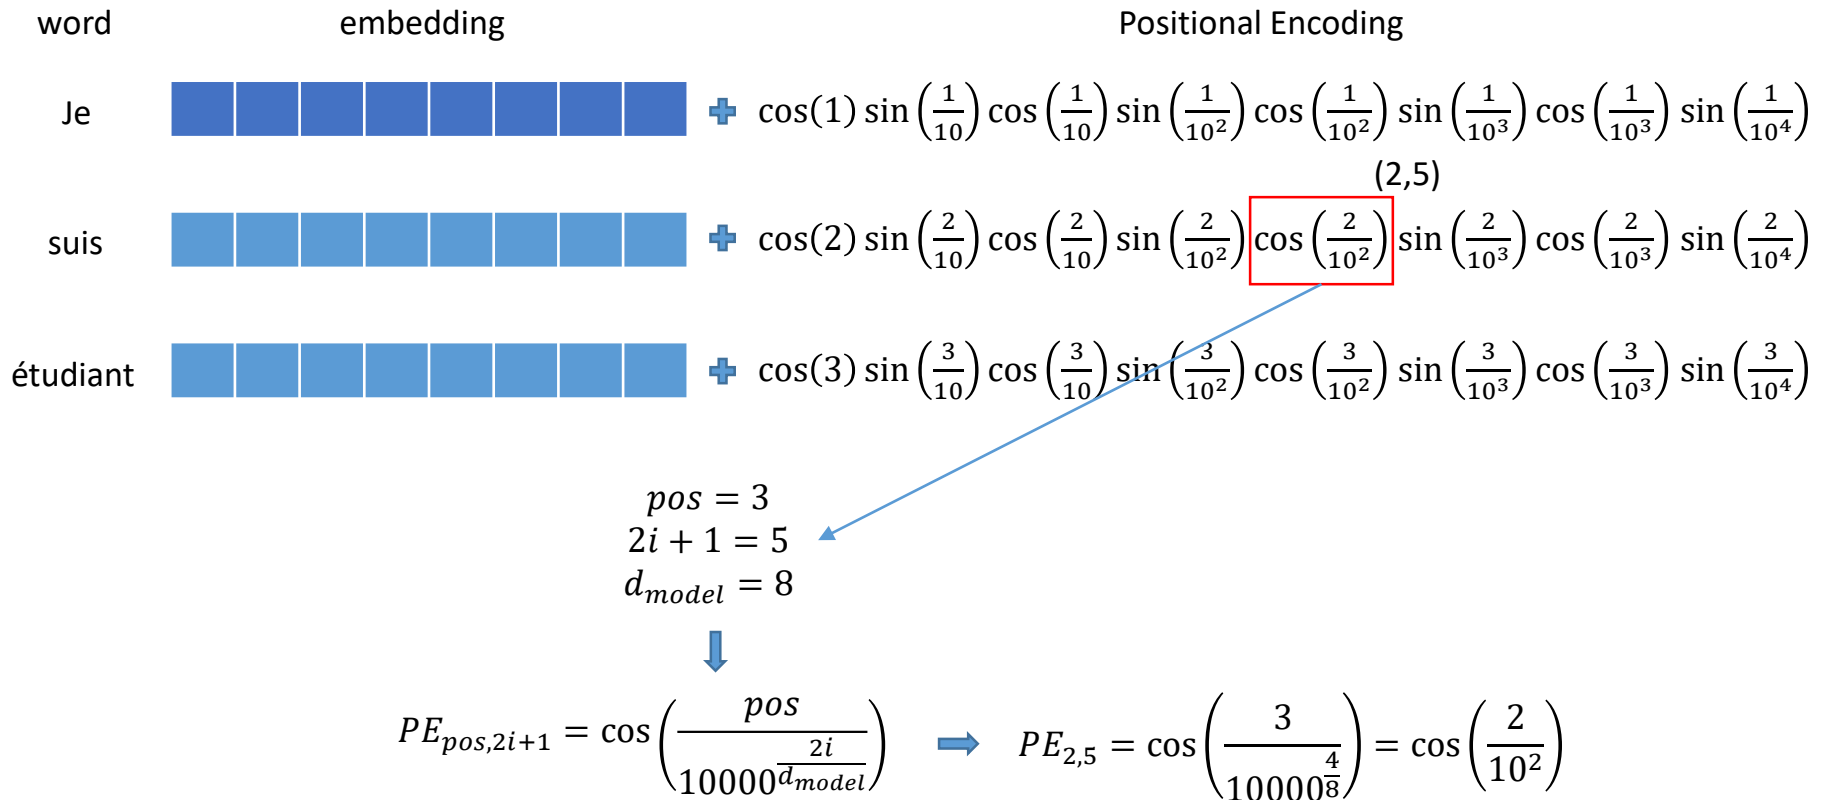
Je




# Positional Encoding



# Positional Encoding




# Positional Encoding






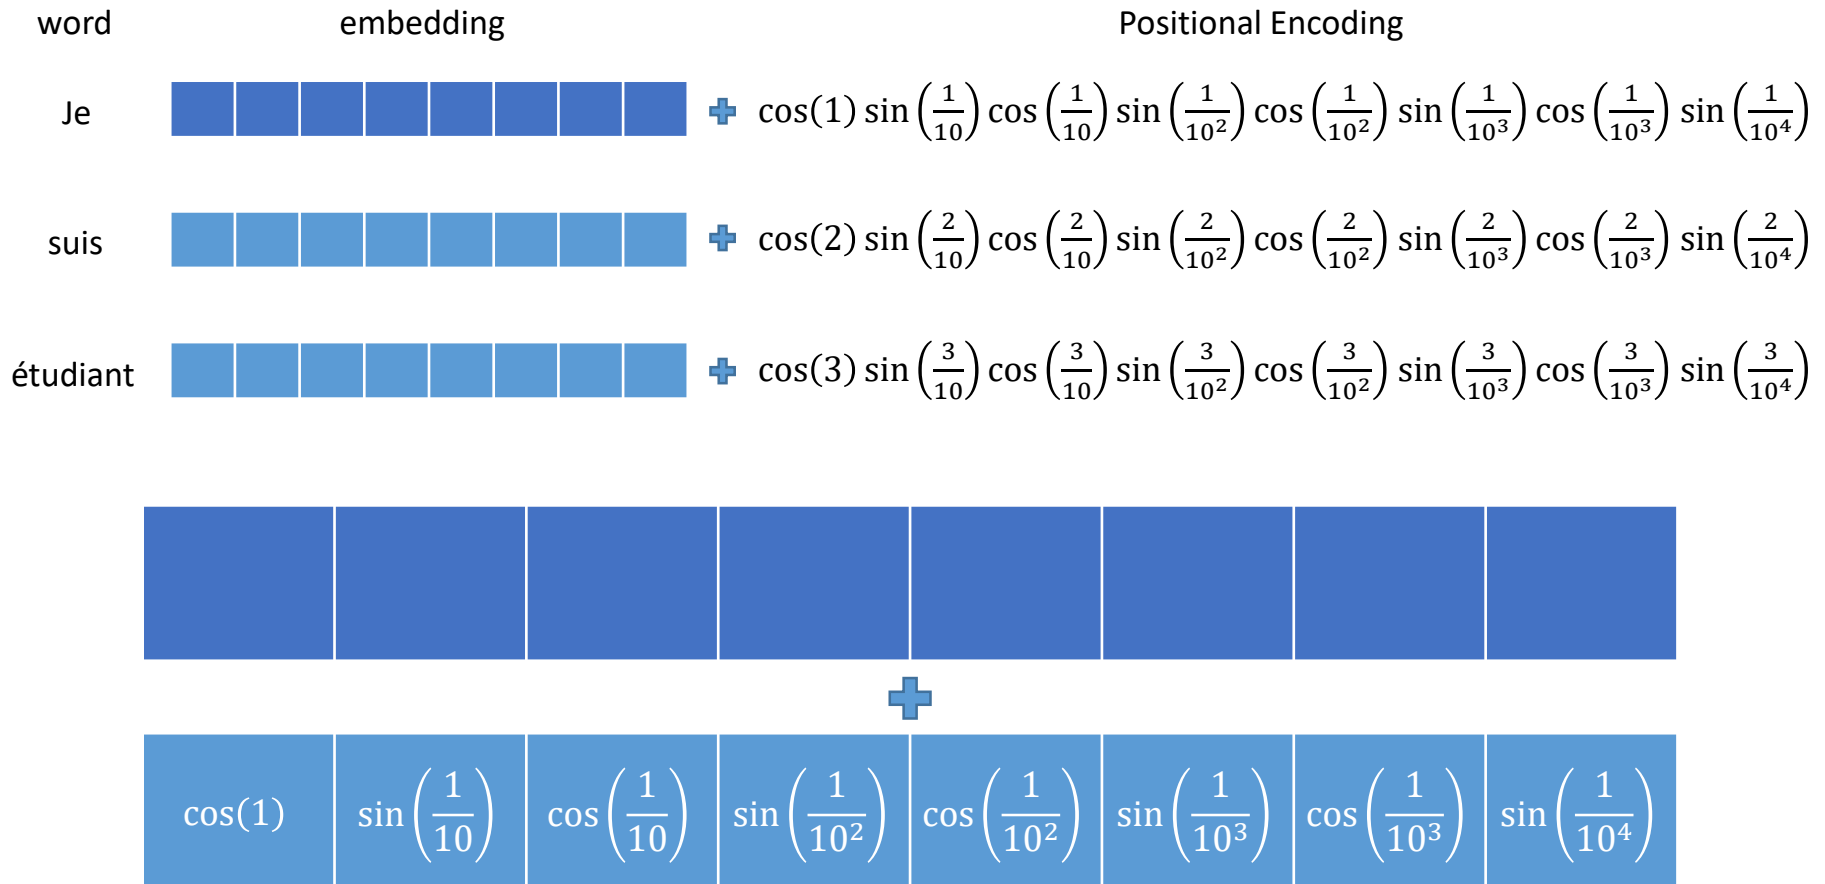
# Positional Encoding



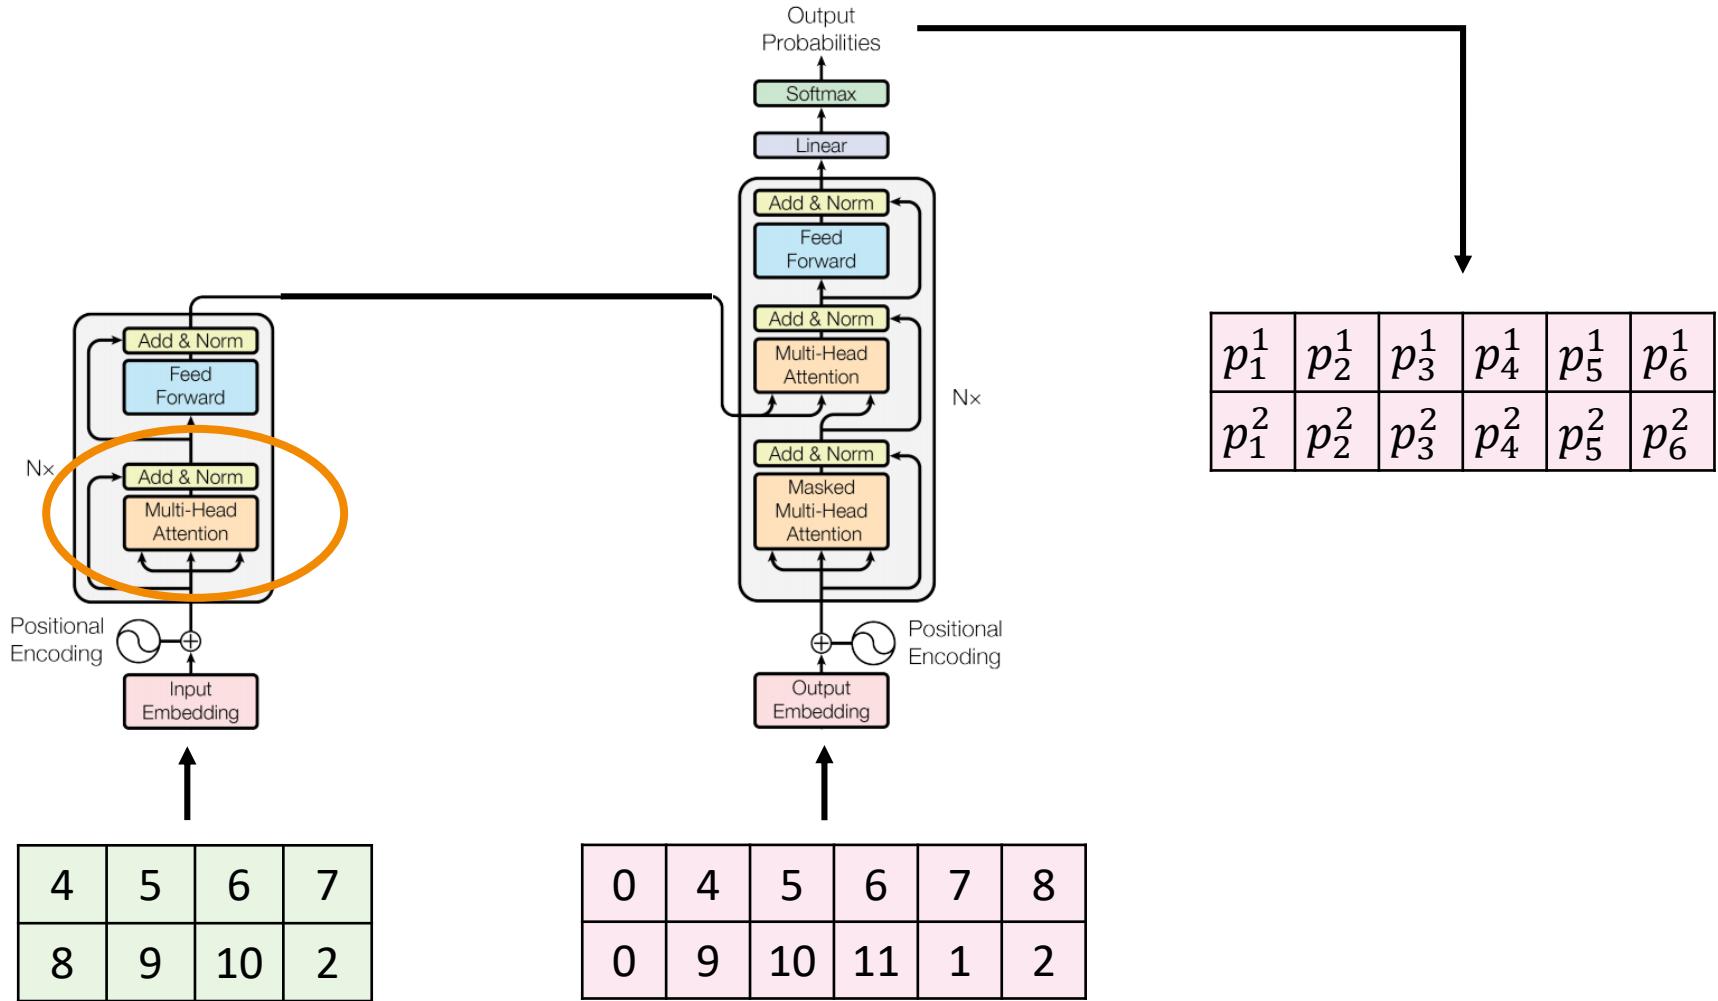
# Positional Encoding



# Positional Encoding



# Multi-Head Attention (Self Attention)



# Multi-Head Attention (Self Attention)

Multi\_Head\_Attention (




0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

)

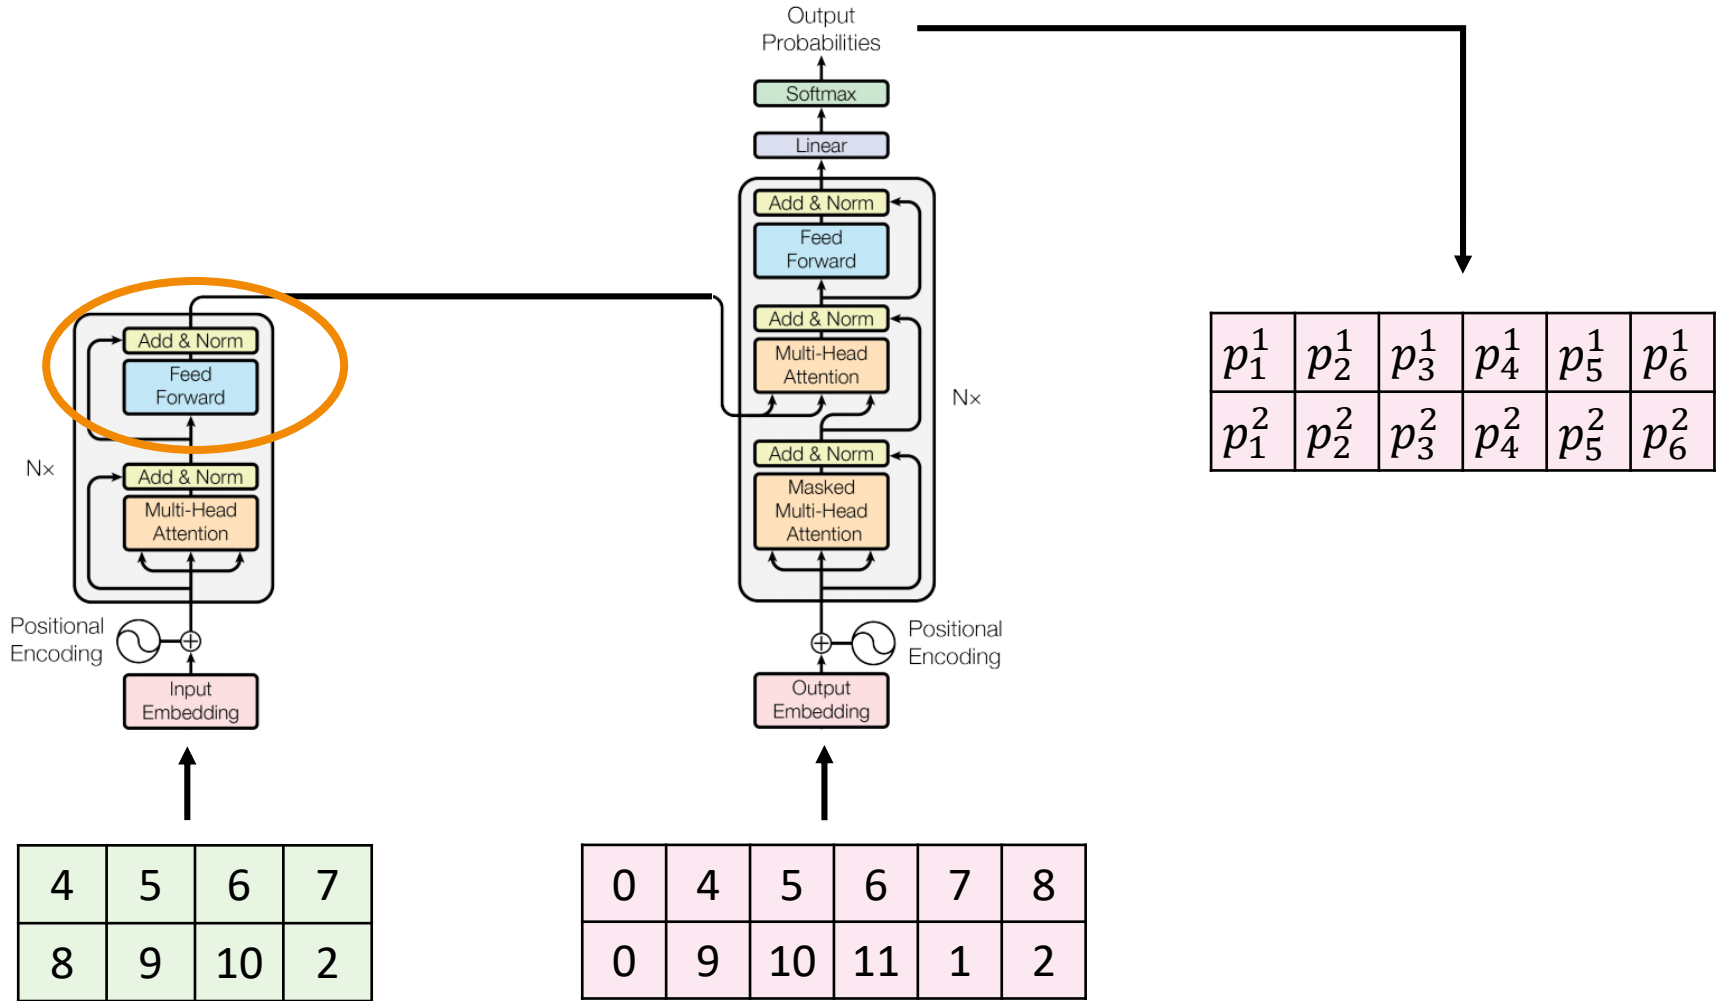
Multi\_Head\_Attention (




0	0	0	1
0	0	0	1
0	0	0	1
0	0	0	1



)

# Position-wise Feed-Forward Networks






# Position-wise Feed-Forward Networks

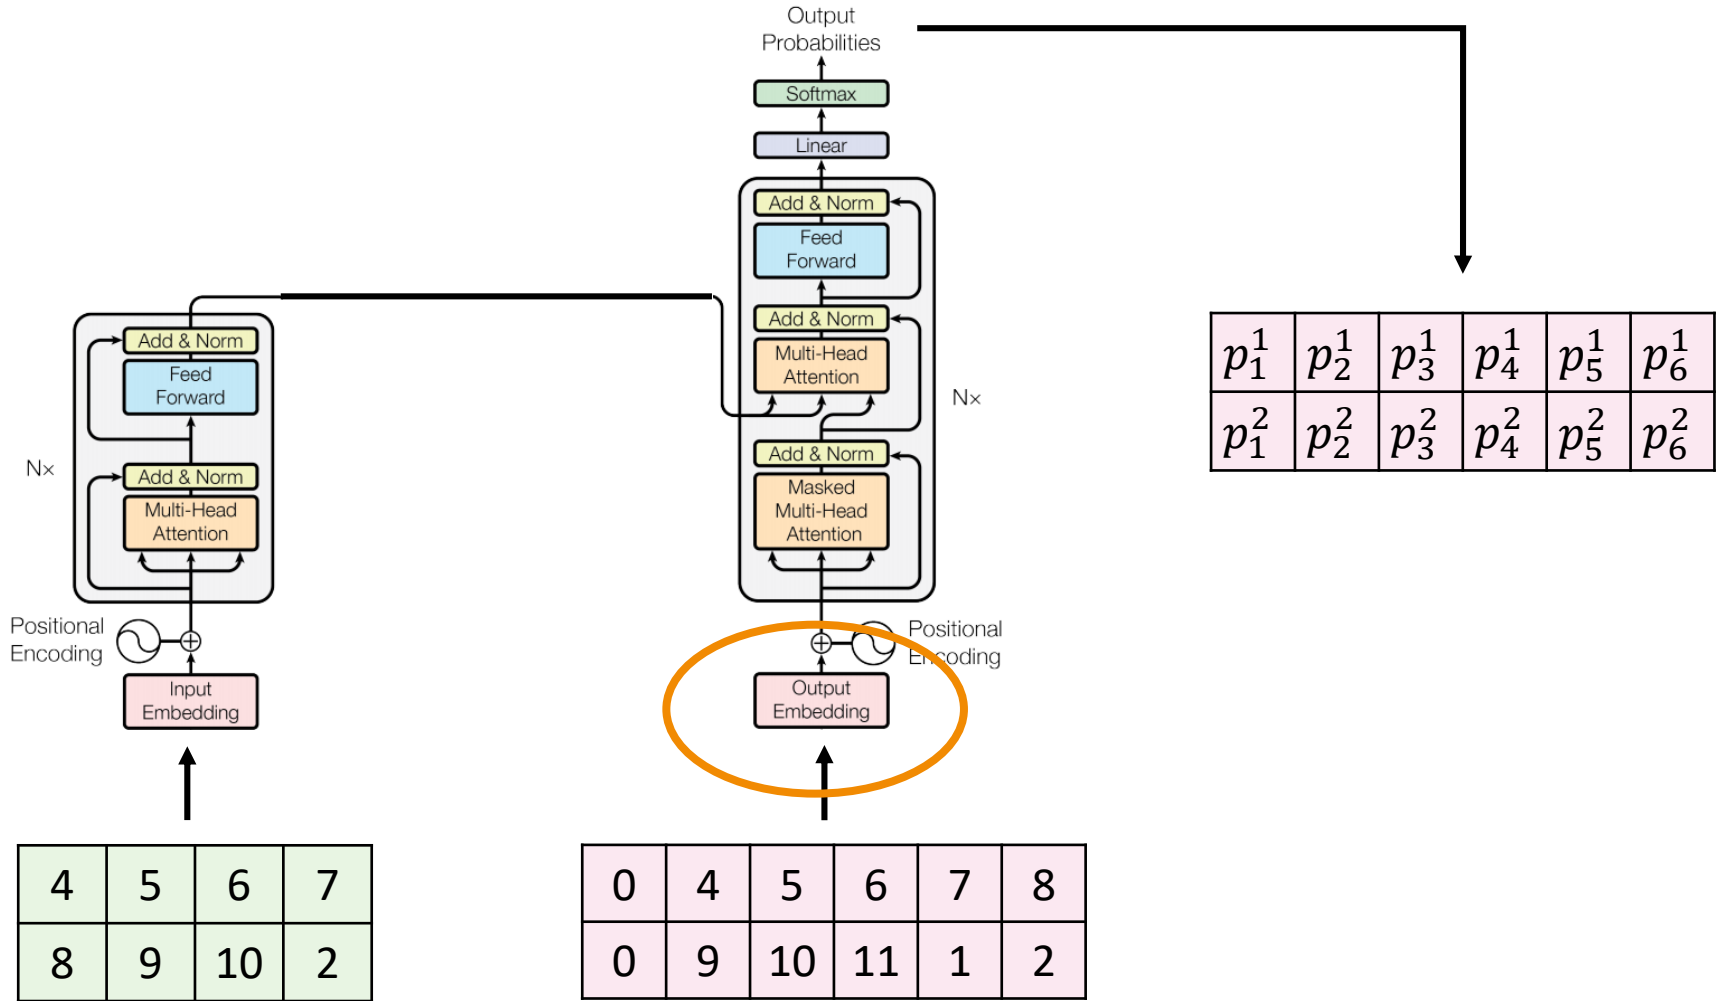
---



Position-wise Feed-Forward Networks

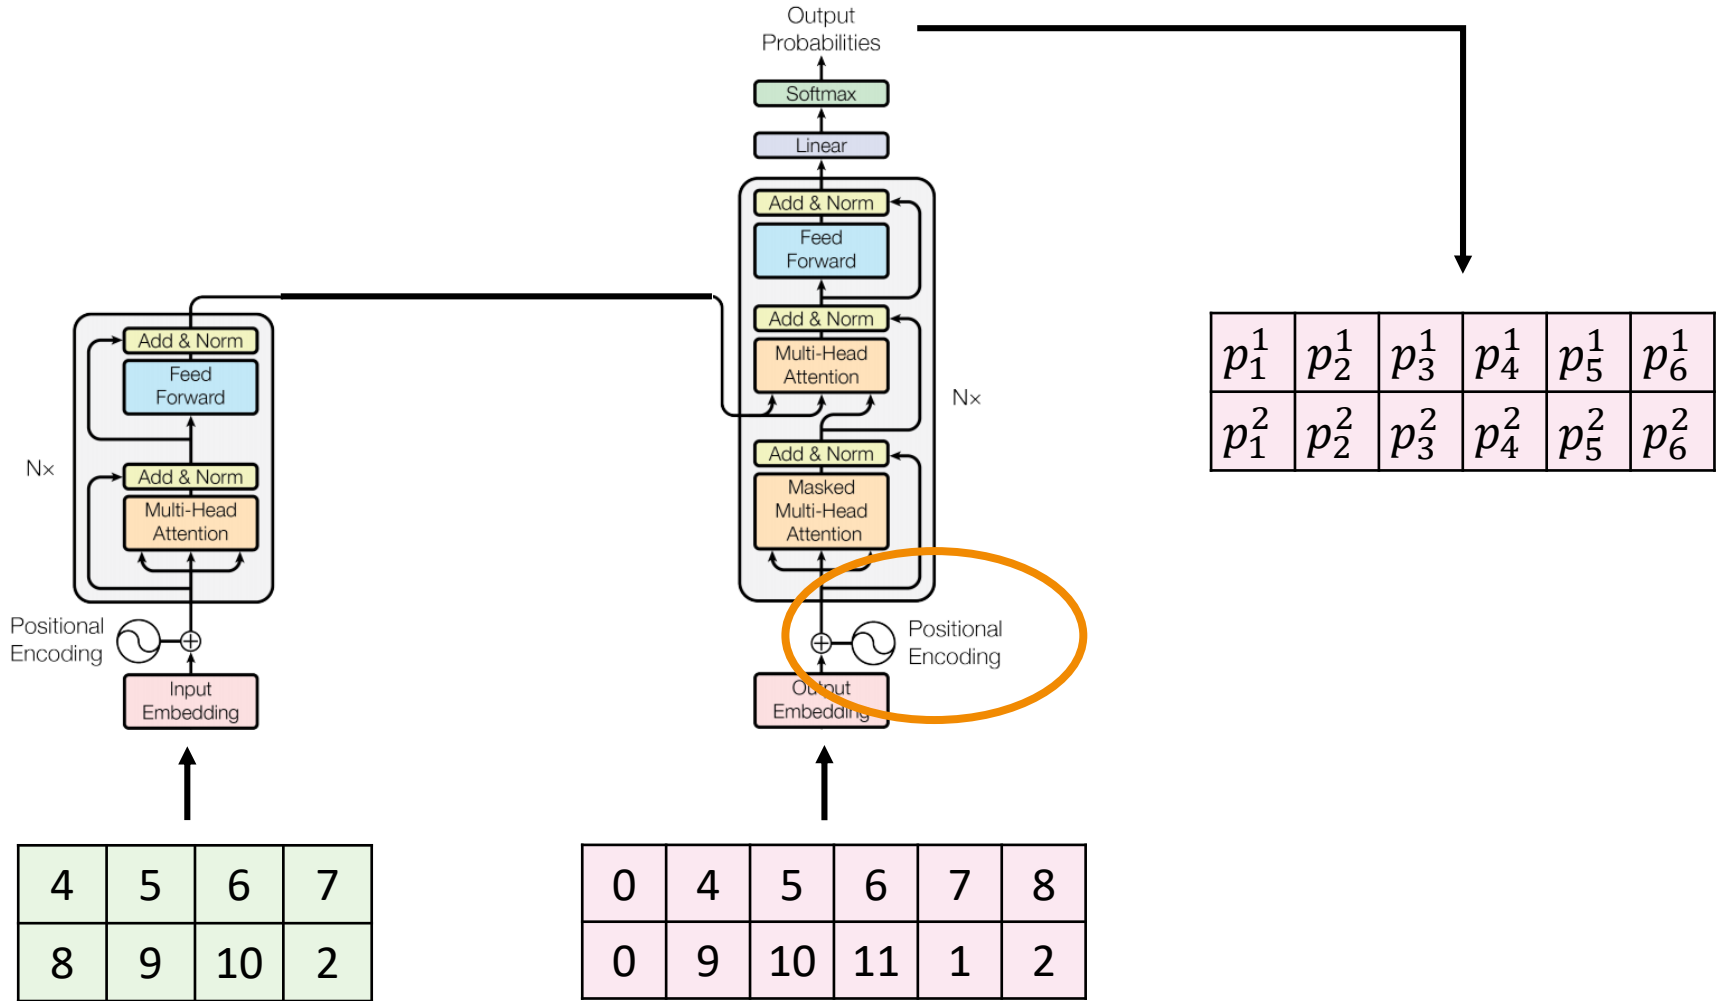


# Output Embedding

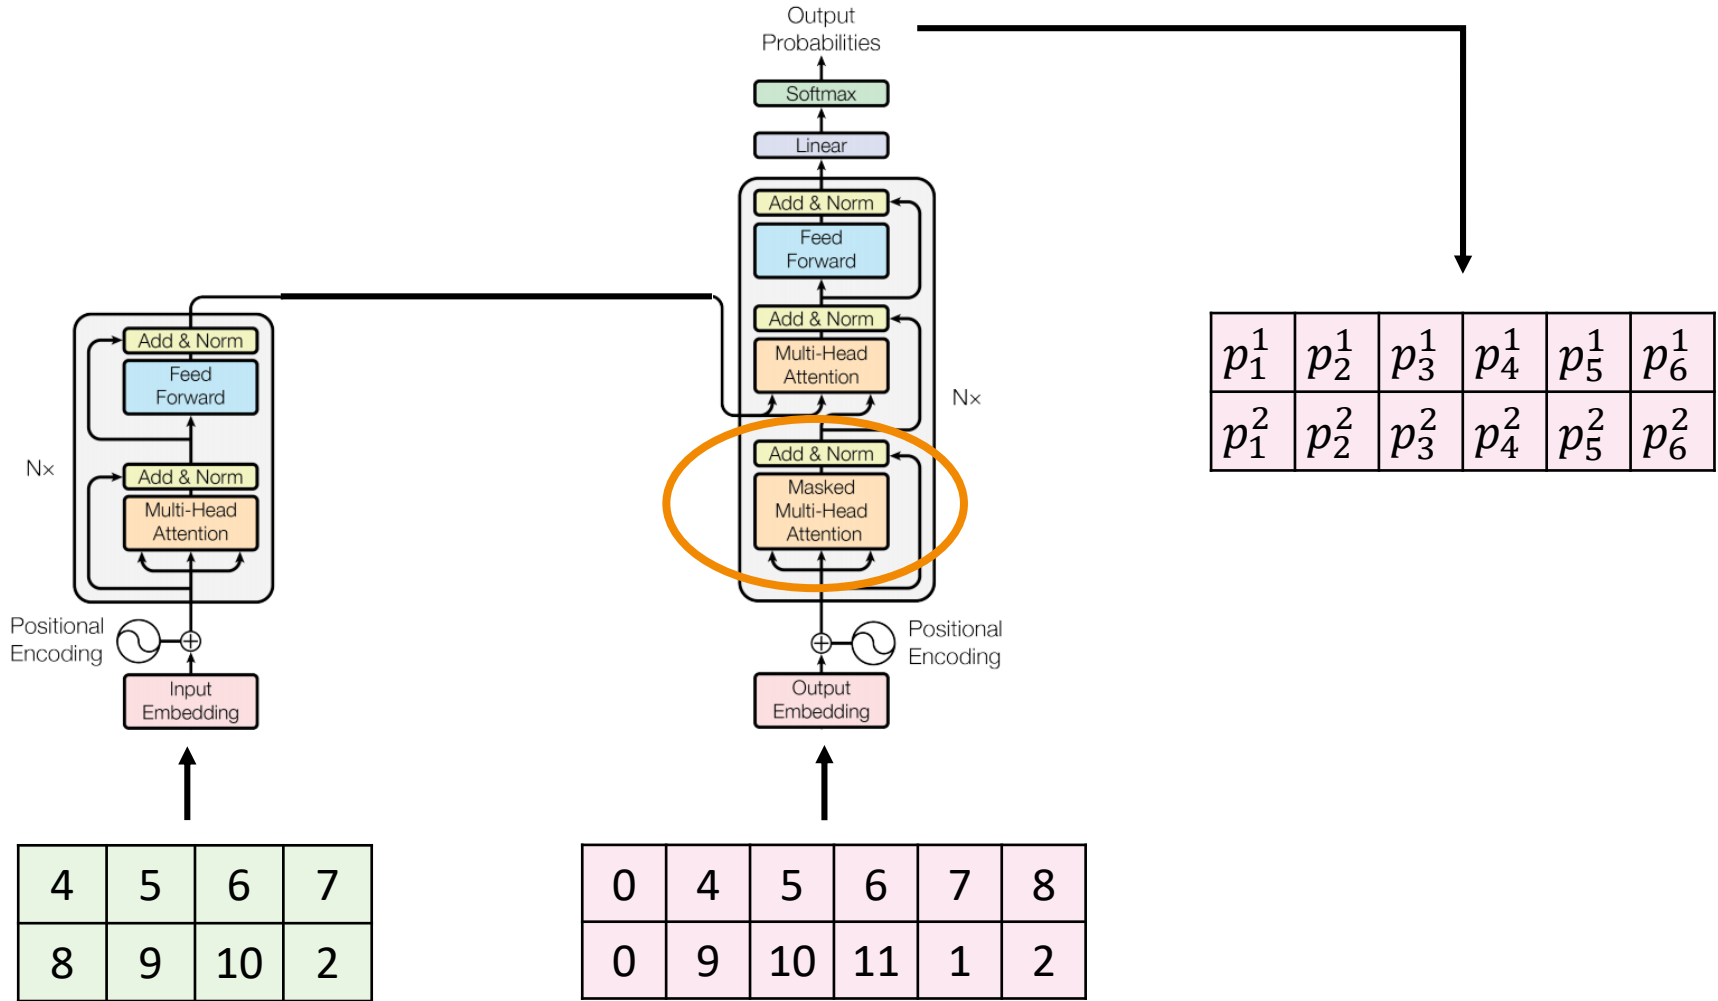




# Positional Encoding



# Multi-Head Attention (Masked Self Attention)



# Multi-Head Attention (Masked Self Attention)

Multi\_Head\_Attention (




0	1	1	1	1	1
0	0	1	1	1	1
0	0	0	1	1	1
0	0	0	0	1	1
0	0	0	0	0	1
0	0	0	0	0	0

)

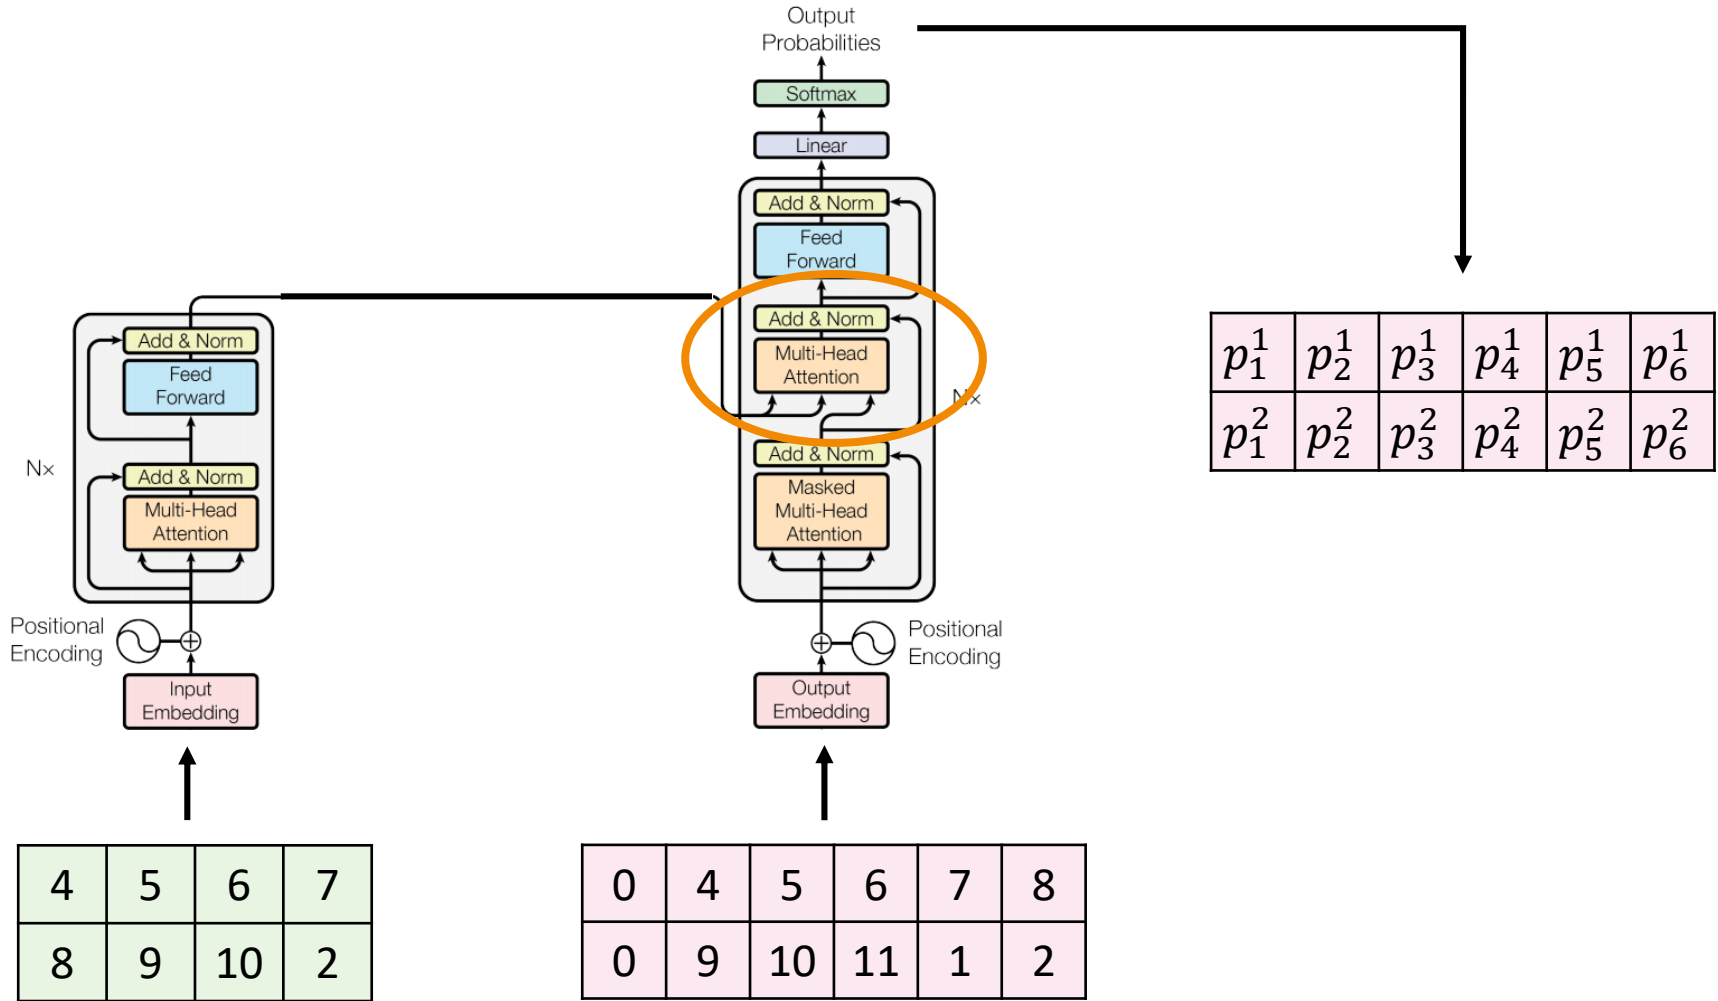
Multi\_Head\_Attention (




0	1	1	1	1	1
0	0	1	1	1	1
0	0	0	1	1	1
0	0	0	0	1	1
0	0	0	0	0	1
0	0	0	0	0	1

)

# Multi-Head Attention (Original Attention)



# Multi-Head Attention (Original Attention)

Multi\_Head\_Attention (




0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

)

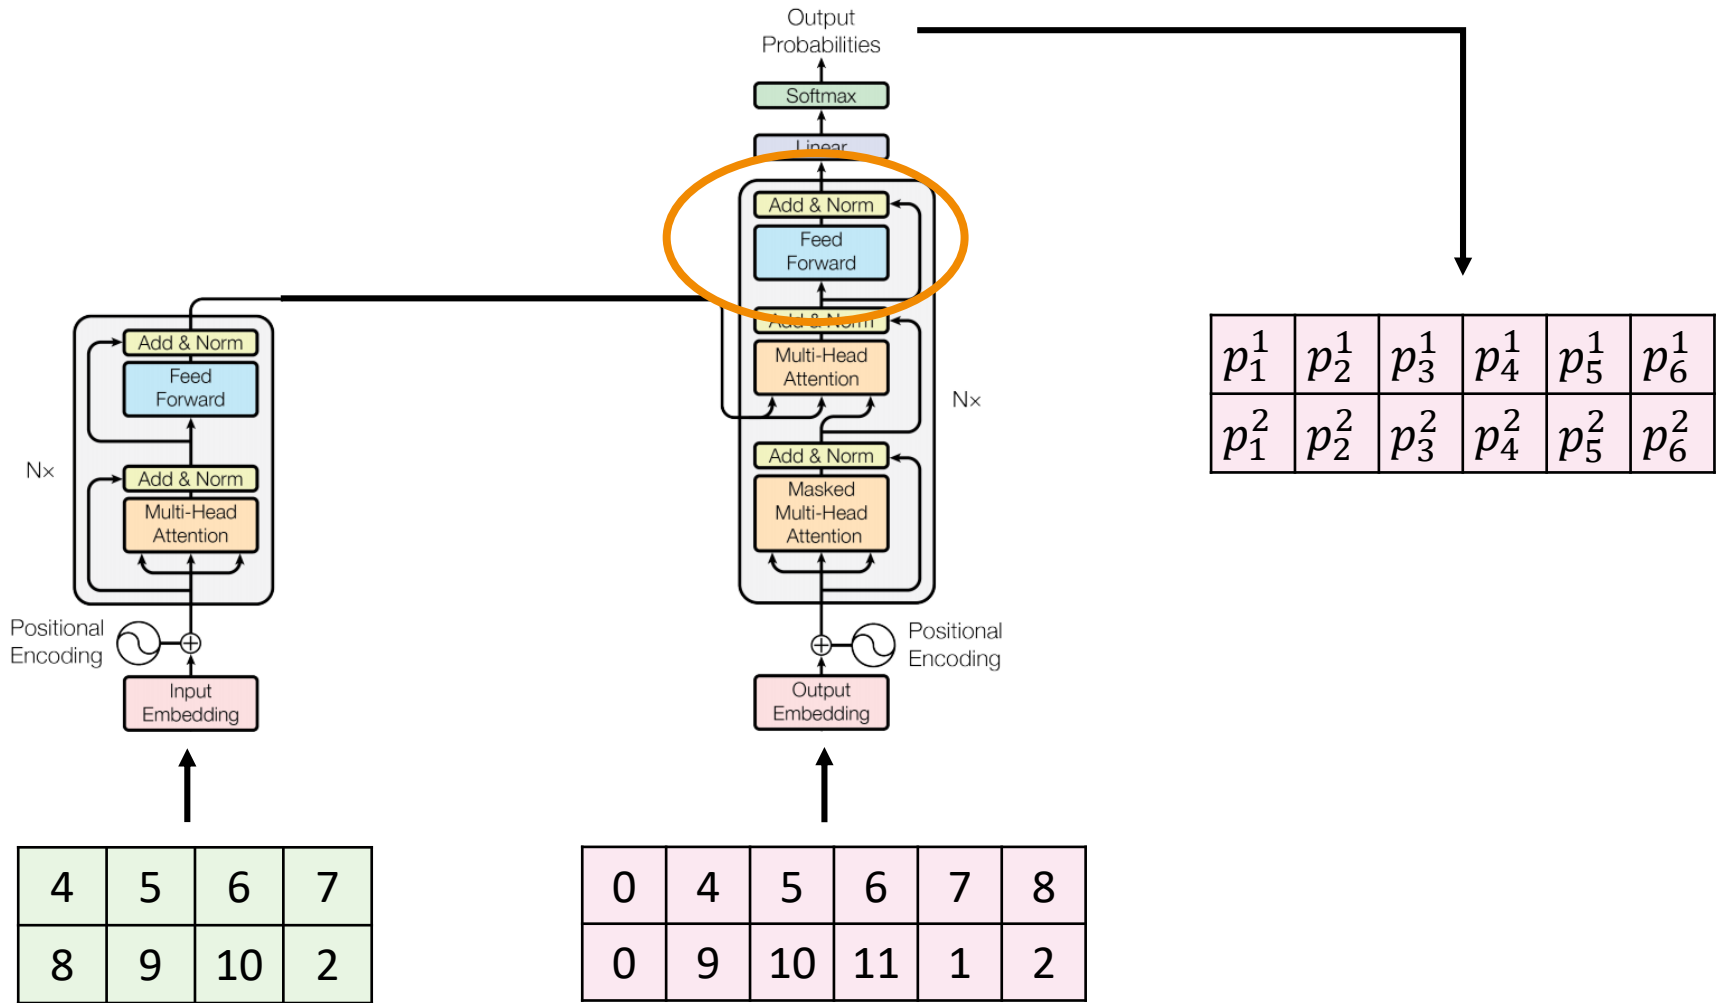
Multi\_Head\_Attention (



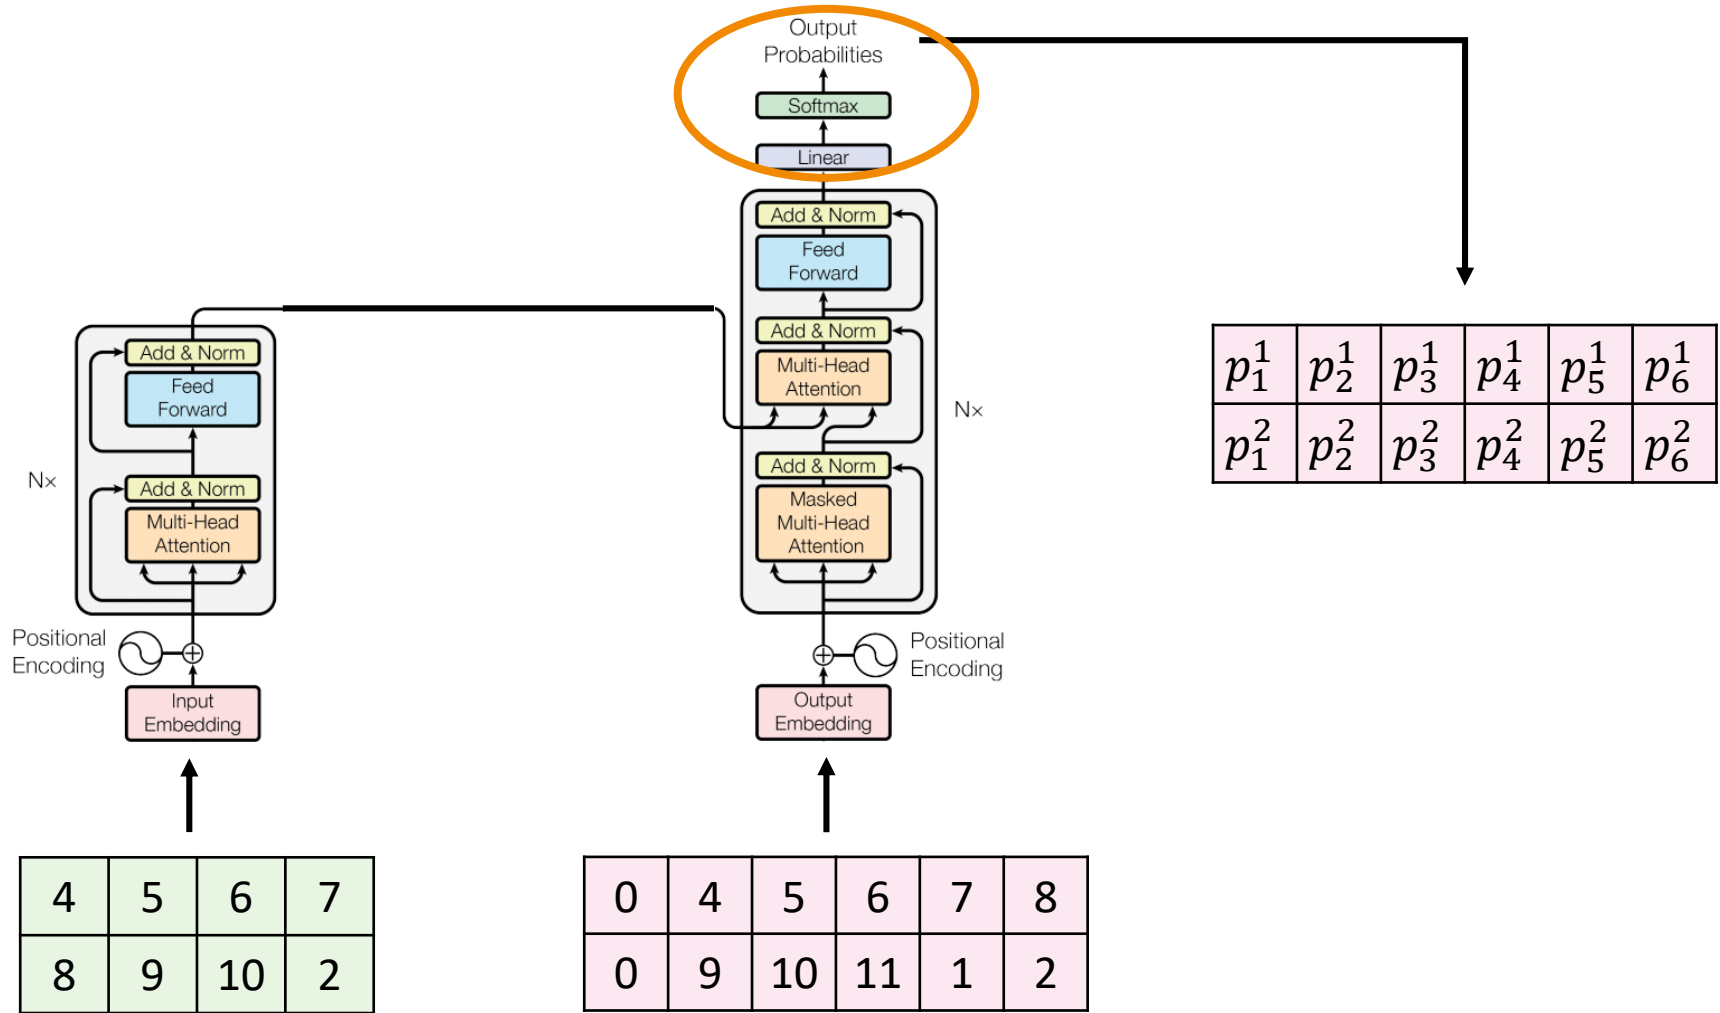

0	0	0	1
0	0	0	1
0	0	0	1
0	0	0	1
0	0	0	1
0	0	0	1

)

# Position-wise Feed-Forward Networks



# Prediction



# References

---

- 1) Attention Is All You Need (NeurIPS'17)
  - <https://papers.nips.cc/paper/7181-attention-is-all-you-need.pdf>
- 2) Neural Machine Translation by Jointly Learning to Align and Translate (ICLR'15)
  - <https://arxiv.org/pdf/1409.0473.pdf>
- 3) Transformers From Scratch
  - <http://www.peterbloem.nl/blog/transformers>
- 4) The Illustrated Transformer
  - <https://jalammar.github.io/illustrated-transformer/>
- 5) The Annotated Transformer
  - <https://nlp.seas.harvard.edu/2018/04/03/attention.html>
- 6) A Brief Overview of Attention Mechanism
  - <https://medium.com/syncedreview/a-brief-overview-of-attention-mechanism-13c578ba9129>



---

# Thank you !

## Any Questions ?